

COMP 737011 - Memory Safety and Programming Language Design

Lecture 3: Heap Attack and Protection

Hui Xu

xuh@fudan.edu.cn



Outline

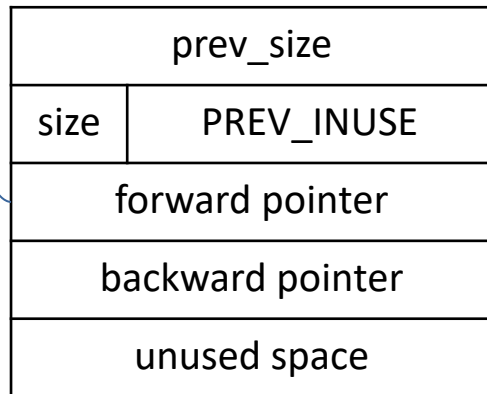
- 1. Heap Analysis
- 2. Heap Attack
- 3. Protection Techniques

1. Heap Analysis


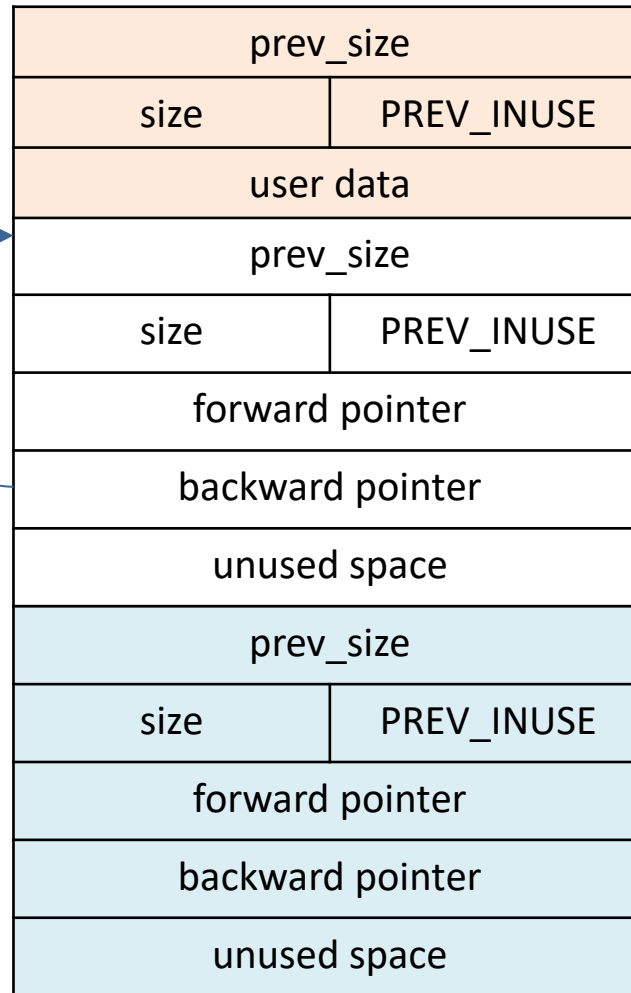
Recall: Chunk Structure

allocated
trunk

free
trunk



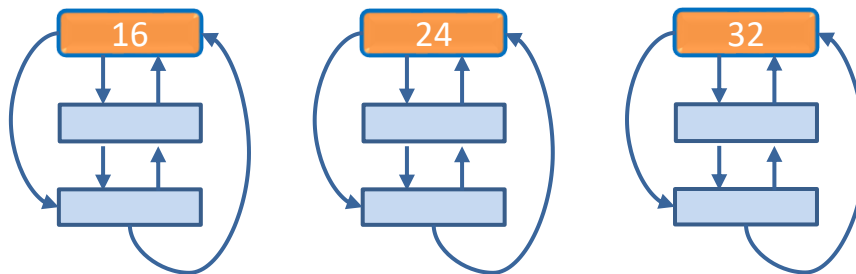
free
trunk



Recall: Doug Lea's Allocator

- Freed memory chunks are managed as bins
 - Regular bins for sizes < 512 bytes are spaced 8 bytes apart
 - Larger bins are approximately logarithmically spaced
- The detailed implementations could vary among allocators

	list	coalesce	data
Fast bin	single-linked	no	small
Regular bin	double-linked	may	could be large



Analyze The Program with GEF

- How many chunks will be allocated?
- What happens to the bins?
- Use the GEF (GDB Enhanced Features) tool for analysis
 - <https://hugsy.github.io/gef/>

```
int main(int argc, char** argv) {  
    char *p[10];  
    for(int i=0; i<10; i++){  
        p[i] = malloc (10 * (i+1));  
        strcpy(p[i], "nowar!!!");  
    }  
  
    for(int i=0; i<10; i++){  
        free(p[i]);  
    }  
    return 0;  
}
```

break 1 →

break 2 →

Disassemble

gef> disass main

Dump of assembler code for function main:

...

0x..1189 <+41>: movsxd rdi,eax

0x..118c <+44>: call 0x1050 <malloc@plt>

0x..1191 <+49>: mov rcx,rax

0x..1194 <+52>: movsxd rax,DWORD PTR [rbp-0x64]

0x..1198 <+56>: mov QWORD PTR [rbp+rax*8-0x60],rcx

0x..119d <+61>: movsxd rax,DWORD PTR [rbp-0x64]

0x..11a1 <+65>: mov rdi,QWORD PTR [rbp+rax*8-0x60]

0x..11a6 <+70>: lea rsi,[rip+0xe57] # 0x2004

0x..11ad <+77>: call 0x1040 <strcpy@plt>

0x..11b2 <+82>: mov eax,DWORD PTR [rbp-0x64]

...

0x..11c0 <+96>: mov DWORD PTR [rbp-0x68],0x0

0x..11c7 <+103>: cmp DWORD PTR [rbp-0x68],0xa

0x..11cb <+107>: jge 0x11ed <main+141>

0x..11d1 <+113>: movsxd rax,DWORD PTR [rbp-0x68]

0x..11d5 <+117>: mov rdi,QWORD PTR [rbp+rax*8-0x60]

0x..11da <+122>: call 0x1030 <free@plt>

...

break 1 →

break 2 →

Check the Allocated Chunk

```
gef> break *main+82
Breakpoint 1 at 0x401191
gef> r
gef> search-pattern nowar
[+] Searching 'nowar' in memory
[+] In '/home/aisr/memory_safety/3-heapattack/a.out'(0x555555556000-0x555555557000), permission=r--
    0x555555556004 - 0x55555555600c → "nowar!!!"
[+] In '/home/aisr/memory_safety/3-heapattack/a.out'(0x555555557000-0x555555558000), permission=r--
    0x555555557004 - 0x55555555700c → "nowar!!!"
[+] In '[heap]'(0x555555559000-0x555555557a000), permission=rw-
    0x5555555592a0 - 0x5555555592a8 → "nowar!!!"
gef> n
gef> search-pattern nowar
...
[+] In '[heap]'(0x555555559000-0x555555557a000), permission=rw-
    0x5555555592a0 - 0x5555555592a8 → "nowar!!!"
    0x5555555592c0 - 0x5555555592c8 → "nowar!!!"
```


Check the Allocated Chunk

```
gef➤ x/30b 0x555555559290
```

```
0x555555559290: 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x00
0x555555559298: 0x21 0x00 0x00 0x00 0x00 0x00 0x00 0x00
0x5555555592a0: 0x6e 0x6f 0x77 0x61 0x72 0x21 0x21 0x21
0x5555555592a8: 0x00 0x00 0x00 0x00 0x00 0x00
```

- chunk size: 0x20
- previous in use: 1

prev_size	
size	PREV_INUSE
data	

- The chunk size is 32 bytes, including the header fields.
- If the previous chunk is in use, the prev_size field can be used to store data of the previous trunk

View The Chunks

```
gef> heap chunks
Chunk(addr=0x55555559010, size=0x290, flags=PREV_INUSE)
  [0x000055555559010    00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00    .....]
Chunk(addr=0x555555592a0, size=0x20, flags=PREV_INUSE)
  [0x0000555555592a0    6e 6f 77 61 72 21 21 21 00 00 00 00 00 00 00 00    nowar!!!.....]
Chunk(addr=0x555555592c0, size=0x20, flags=PREV_INUSE)
  [0x0000555555592c0    6e 6f 77 61 72 21 21 21 00 00 00 00 00 00 00 00    nowar!!!.....]
Chunk(addr=0x555555592e0, size=0x20d30, flags=PREV_INUSE)
  [0x0000555555592e0    00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00    .....]
Chunk(addr=0x555555592e0, size=0x20d30, flags=PREV_INUSE) ← top chunk
```

- The chunk sizes are both 0x20 for the first two malloc
- 16 bytes spaced apart

After Several Iterations

```
gef> heap chunks
Chunk(addr=0x55555559010, size=0x290, flags=PREV_INUSE)
  [0x000055555559010  00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00  .....]
Chunk(addr=0x555555592a0, size=0x20, flags=PREV_INUSE)
  [0x0000555555592a0  6e 6f 77 61 72 21 21 21 00 00 00 00 00 00 00 00  nowar!!!.....]
Chunk(addr=0x555555592c0, size=0x20, flags=PREV_INUSE)
  [0x0000555555592c0  6e 6f 77 61 72 21 21 21 00 00 00 00 00 00 00 00  nowar!!!.....]
Chunk(addr=0x555555592e0, size=0x30, flags=PREV_INUSE)
  [0x0000555555592e0  6e 6f 77 61 72 21 21 21 00 00 00 00 00 00 00 00  nowar!!!.....]
Chunk(addr=0x55555559310, size=0x30, flags=PREV_INUSE)
  [0x000055555559310  6e 6f 77 61 72 21 21 21 00 00 00 00 00 00 00 00  nowar!!!.....]
Chunk(addr=0x55555559340, size=0x40, flags=PREV_INUSE)
  [0x000055555559340  6e 6f 77 61 72 21 21 21 00 00 00 00 00 00 00 00  nowar!!!.....]
Chunk(addr=0x55555559380, size=0x50, flags=PREV_INUSE)
  [0x000055555559380  6e 6f 77 61 72 21 21 21 00 00 00 00 00 00 00 00  nowar!!!.....]
Chunk(addr=0x555555593d0, size=0x50, flags=PREV_INUSE)
  [0x0000555555593d0  6e 6f 77 61 72 21 21 21 00 00 00 00 00 00 00 00  nowar!!!.....]
Chunk(addr=0x55555559420, size=0x60, flags=PREV_INUSE)
  [0x000055555559420  6e 6f 77 61 72 21 21 21 00 00 00 00 00 00 00 00  nowar!!!.....]
Chunk(addr=0x55555559480, size=0x70, flags=PREV_INUSE)
  [0x000055555559480  6e 6f 77 61 72 21 21 21 00 00 00 00 00 00 00 00  nowar!!!.....]
Chunk(addr=0x555555594f0, size=0x70, flags=PREV_INUSE)
  [0x0000555555594f0  6e 6f 77 61 72 21 21 21 00 00 00 00 00 00 00 00  nowar!!!.....]
Chunk(addr=0x55555559560, size=0x20ab0, flags=PREV_INUSE)
  [0x000055555559560  00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00  .....]
Chunk(addr=0x55555559560, size=0x20ab0, flags=PREV_INUSE) ← top chunk
```

View The Bins (tcachebins)

```
gef> heap bins
_____ Tcachebins for thread 1
_____

All tcachebins are empty
_____ Fastbins for arena at 0x7ffff7facb80
_____

Fastbins[idx=0, size=0x20] 0x00
Fastbins[idx=1, size=0x30] 0x00
Fastbins[idx=2, size=0x40] 0x00
Fastbins[idx=3, size=0x50] 0x00
Fastbins[idx=4, size=0x60] 0x00
Fastbins[idx=5, size=0x70] 0x00
Fastbins[idx=6, size=0x80] 0x00
_____ Unsorted Bin for arena at 0x7ffff7facb80
_____

[+] Found 0 chunks in unsorted bin.
_____ Small Bins for arena at 0x7ffff7facb80
_____

[+] Found 0 chunks in 0 small non-empty bins.
_____ Large Bins for arena at 0x7ffff7facb80
_____

[+] Found 0 chunks in 0 large non-empty bins.
```

- Freed chunks will be added to tcachebins (new in libc 2.6)

View The Bins (tcachebins)

- Freed chunks after several iterations.

```
gef> heap bins
```

```
———— Tcachebins for thread 1 —————  
Tcachebins[idx=0, size=0x20, count=2]  
← Chunk(addr=0x5555555592c0, size=0x20, flags=PREV_INUSE)  
← Chunk(addr=0x5555555592a0, size=0x20, flags=PREV_INUSE)  
Tcachebins[idx=1, size=0x30, count=2]  
← Chunk(addr=0x555555559310, size=0x30, flags=PREV_INUSE)  
← Chunk(addr=0x5555555592e0, size=0x30, flags=PREV_INUSE)  
Tcachebins[idx=2, size=0x40, count=1]  
← Chunk(addr=0x555555559340, size=0x40, flags=PREV_INUSE)  
Tcachebins[idx=3, size=0x50, count=2]  
← Chunk(addr=0x5555555593d0, size=0x50, flags=PREV_INUSE)  
← Chunk(addr=0x555555559380, size=0x50, flags=PREV_INUSE)  
Tcachebins[idx=4, size=0x60, count=1]  
← Chunk(addr=0x555555559420, size=0x60, flags=PREV_INUSE)
```

prev_size	
size	PREV_INUSE
forward pointer (data)	
backward pointer (optional)	
unused space	

Summarization of Allocation Behaviors

- The first malloc reserves a large chunk (32KB)
 - The first 0x290 bytes used for bin management
 - The following mallocs obtain trunks from the reserved trunk.
- Freed chunks are added to tcachebins
 - Single-linked list, first-in-last-out
 - Max length of the list in each bin: 7
- Exceeding chunks will be put into fastbins

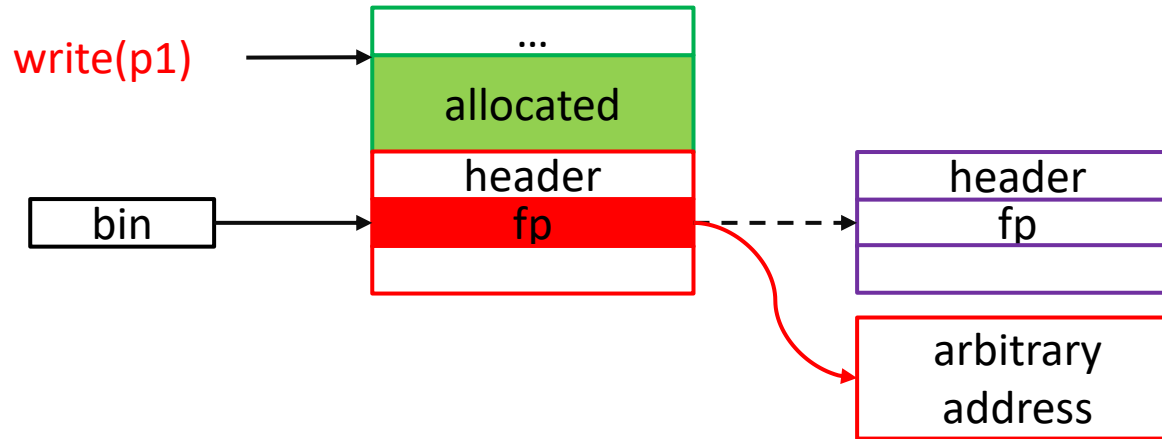
2. Heap Attack

Heap Vulnerabilities

- Heap overflow
- Use after free
- Double free

Heap Overflow

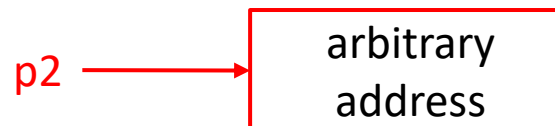
Step1: modify the fp of the next chunk to an arbitrary address



Step2: allocate the next chunk via malloc()

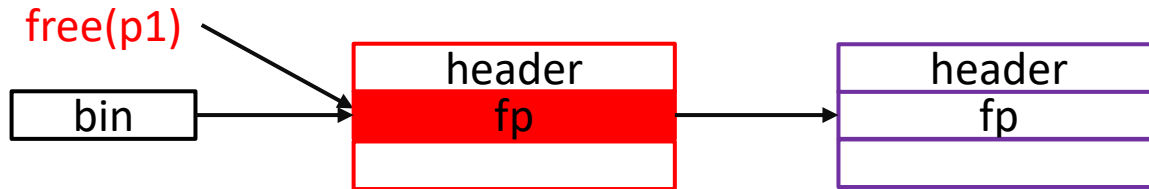


Step3: call malloc() again

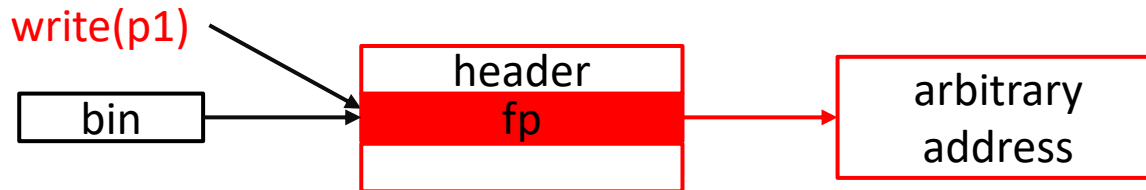


Use After Free

Step1: free(p1)



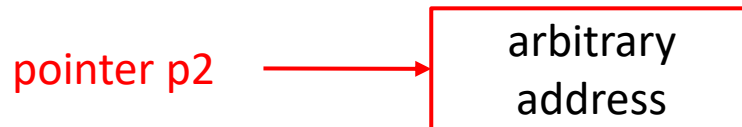
Step2: modify fp to an arbitrary address



Step3: malloc()




Step4: malloc() again to obtain a pointer to the arbitrary address




Double Free

```
//free(new):  
first = header->next  
header->next = new  
new->next = first
```


Step1: free(p1)



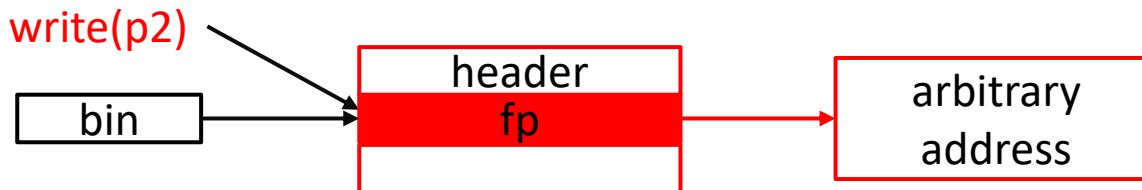
Step2: free(p1) again




Step3: call malloc()



Step4: modify fp to an arbitrary address



Step5: malloc() twice to obtain a pointer to the arbitrary address



Address of Attacking Interest

- Return Address:
 - similar as buffer overflow
- Global Offset Table (GOT):
 - a table for dynamic linkage or position-independent code
 - change the table entries, *e.g.*, address of strcpy()
- Virtual Method Table (vtable):
 - abstract functions of C++/Rust

3. Protection Techniques

Detect Bugs in Allocator?

- Use static analysis or dynamic analysis?
- Detect invalid behaviors during malloc/free?
 - Chunk addresses should within the valid range?
 - A free chunk should not be freed again?
 - More fine-grained strategies?
- Detect invalid behaviors during read/write?
 - Overhead issues
- Increase the difficulty of heap attack?

Static Analysis Is Hard

- The fundamental point-to/alias analysis is NP-hard
- Several typical performance issues to consider
 - Flow-sensitivity: consider the order of statements
 - Path-sensitivity: analyze the result for each path
 - Context-sensitivity: inter-procedural issues
 - Field-sensitivity: how to model the members of objects
- Related papers:
 - Lee, *et al.* "Preventing Use-after-free with Dangling Pointers Nullification." NDSS 2015.
 - Van Der Kouwe, *et al.* "Dangsan: Scalable use-after-free detection." EuroSys 2017.
- We will have a class for the topic

Dynamic Approach Is Expensive

- Runtime detection mechanisms are needed
 - E.g., offset could be used => boundary check
- Trade-off between security and efficiency
- Mechanisms used in current allocators
 - alignment check
 - fasttop
 - canary

Alignment Check: Invalid Pointer Detection

- The following code is used within the function `_int_free()`
- Free a misaligned chunk is invalid

```
#define CHUNK_HDR_SZ (2 * SIZE_SZ) // 2 * size_t, 16 byte in x86-64
#define MALLOC_ALIGN_MASK (MALLOC_ALIGNMENT - 1)
#define misaligned_chunk(p) \
    ((uintptr_t)(MALLOC_ALIGNMENT == CHUNK_HDR_SZ ? \
        (p) : chunk2mem (p)) & MALLOC_ALIGN_MASK)

/* Little security check which won't hurt performance: the
   allocator never wraps around at the end of the address space.
   Therefore we can exclude some size values which might appear
   here by accident or by "design" from some intruder. */
if (__builtin_expect ((uintptr_t) p > (uintptr_t) -size, 0)
    || __builtin_expect (misaligned_chunk (p), 0))
    malloc_printerr ("free(): invalid pointer");
```

Fasttop: Double Free Detection

- Fasttop: pointer address should not be just freed
- Also used in the function of `_int_free()`

```
unsigned int idx = fastbin_index(size);  
mfastbinptr fb = &fastbin (av, idx); //av is the malloc_state  
mchunkptr old = *fb;  
if (__builtin_expect (old == p, 0))  
    malloc_printerr ("double free or corruption (fasttop)");
```

Canary (tcache_key): Double Free Detection

- Used only when USE_TCACHE is enabled
- Call tcache_put() in _init_malloc() to store the key

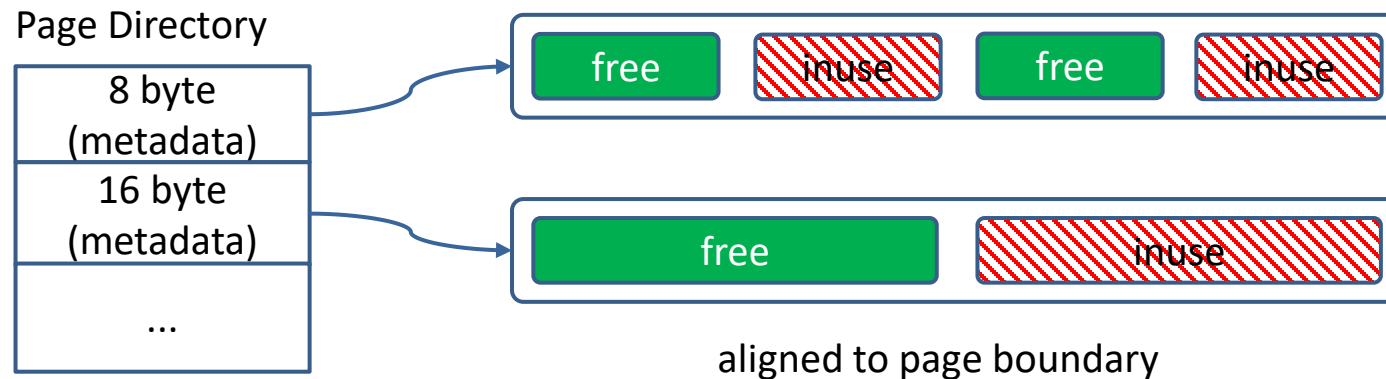
```
typedef struct tcache_entry {  
    struct tcache_entry *next;  
    uintptr_t key; //double free flag  
} tcache_entry;  
  
tcache_put (mchunkptr chunk, size_t tc_idx){  
    tcache_entry *e = (tcache_entry *) chunk2mem (chunk);  
    e->key = tcache_key;  
    ...  
}
```

- Check if content is still the key in the function of _int_free()

```
if (__glibc_unlikely (e->key == tcache_key)) {  
    ...//probe the issue  
}
```

More Approaches: BiBOP-Style Heap

- Big Bag of Pages:
 - contiguous areas of a multiple page size
 - each page has the same sized chunks
 - store heap metadata out-of-band (more secure)
- Originally proposed in PHKmalloc (OpenBSD)



More Papers to Read

- Berger, et al. "DieHard, Probabilistic memory safety for unsafe languages." *PLDI*, 2006.
- Novark, et al. "DieHarder: securing the heap." *CCS*, 2010.
- Akritidis. "Cling: A memory allocator to mitigate dangling pointers." *USENIX Security*, 2010.
- Sam, *et al.* "Freeguard: A faster secure heap allocator." *CCS*, 2017.

Programming Language Design

- Rust ownership-based mechanism
 - prohibit shared mutable aliases
 - no dangling pointer => preventing use after free, double free
- Shared mutable aliases should be wrapped with RC type
 - similar to `shared_ptr` in C++
- We will have a class for the topic

In Class Practice

- Write a C program with one of the following bugs and show how you can manipulate the free list with the bug.
 - Heap overflow
 - Use after free
 - Double free
- Hint:
 - Use the GEP tool to probe the trunks
 - You may encounter some detection techniques for double free

Solution

Solution: Use After Free

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

void main(void)
{
    char* p1 = malloc (22);
    char* p2 = malloc (22);
    free(p2);
    free(p1);
    *(int *) p1 = 0x411112;
    p1 = malloc(22);
    p2 = malloc(22);
    printf("Allocated memory address: %x\n", p2);
}
```

Solution: Double Free

```
void main(void)
{
    char* p1 = malloc (22);
    free(p1);
    p1[9] = 0x0; //overwrite e-key for double check
    free(p1);
    *(int *) p1 = 0x411112;
    p1 = malloc(22);
    p1 = malloc(22);
    printf("Allocated memory address: %x\n", p1);
}
```