

# Lecture 4: Memory Exhaustion

Xu, Hui

[xuh@fudan.edu.cn](mailto:xuh@fudan.edu.cn)



# Outline

- 1. Stack Overflow
- 2. Heap Exhaustion
- 3. Exception Handling
- 4. Stack Unwinding

# 1. Stack Overflow

---

# Warm Up

- Can you find a list to overflow the stack?

```
struct List{  
    int val;  
    struct List* next;  
};
```


```
void process(struct List* l, int cnt){  
    printf("%d\n", cnt);  
    if(l->next != NULL)  
        process(l->next, ++cnt);  
}
```

- Sample solution

```
void main(void){  
    sethandler(handler);  
    struct List* list = malloc(sizeof(struct List));  
    list->val = 1;  
    list->next = list;  
    process(list, 0);  
}
```

# Stack Size is Limited

- Default stack size: 8MB for each thread in Linux
  - You may check the setting with the ulimit command
- Reaching the limit would cause stack overflow
- Why not use a large stack?
  - Mainly used to save the contexts of function calls
  - Developers should not place large data on stack
- Vulnerable code: recursive function calls



```
#: ulimit -a
max locked memory      (kbytes, -l) 65536
max memory size        (kbytes, -m) unlimited
open files             (-n) 1024
pipe size              (512 bytes, -p) 8
POSIX message queues   (bytes, -q) 819200
stack size             (kbytes, -s) 8192
max user processes     (-u) 30687
```

# You May Change The Stack Limit

- System users: ulimit command

```
aisr@aisr:~$ ulimit -s unlimited
aisr@aisr:~$ ulimit -a
max locked memory      (kbytes, -l) 65536
max memory size        (kbytes, -m) unlimited
open files             (-n) 1024
pipe size              (512 bytes, -p) 8
POSIX message queues   (bytes, -q) 819200
stack size             (kbytes, -s) unlimited
```

Set the stack size to unlimited

- Developers: use the setrlimit() function

```
struct rlimit r;
int result;
result = getrlimit(RLIMIT_STACK, &r);
fprintf(stderr, "stack result = %d\n", r.rlim_cur);
r.rlim_cur = 64 * 1024L * 1024L;
result = setrlimit(RLIMIT_STACK, &r);
result = getrlimit(RLIMIT_STACK, &r);
fprintf(stderr, "stack result = %d\n", r.rlim_cur);
```

Set the stack size to 64 MB

# How to Handle Stack Overflow?

- The OS usually kills the process directly. Why?
- We can register a handler for the SIGSEGV signal
- Executing the event handler needs an extra stack
  - need to register another stack with enough space
- You will learn this in your in-class practice

## 2. Heap Exhaustion

---



# Overcommit

- A lazy mode memory allocation mechanism
  - malloc() successful does not mean the physical memory is allocated
  - The physical memory is allocated when being accessed
- Linux has three options
  - 1: always overcommit, never check
  - 2: always check, never overcommit
  - 0: heuristic overcommit (this is the default)

```
#: sudo sysctl -w vm.overcommit_memory=2
```

# Overcommit: Example

- Try the following program with different settings

```
#define LARGE_SIZE 1024L*1024L*1024L*256L
```

→ 256 GB

```
void main(void){
```

```
    char* p = malloc (LARGE_SIZE);
```

```
    if(p == 0) {
```

→ allocation failure

```
        printf("malloc failed\n");
```

```
    } else {
```

→ allocation successful  
access the memory

```
        memset (p, 1, LARGE_SIZE);
```

```
    }
```

```
}
```

```
#: sudo sysctl -w vm.overcommit_memory=1
```

```
#:~/4-memoxhaustion$ ./a.out
```

```
Killed
```

→ killed by the OS

```
#: sudo sysctl -w vm.overcommit_memory=2
```

```
#:~/4-memoxhaustion$ ./a.out
```

```
malloc failed
```

→ allocation failure

# How to Handle Heap Exhaustion?

- Always check: based on the return value of `malloc()`
  - returns 0 if fails
- Overcommit: could be killed by the OS
  - register a handler for the SIGKILL signal?
- Too Small to Fail & OOM Killer
  - If the required space is small ( $< 8$  pages), `malloc()` should never fail when overcommit is enabled
  - If no enough memory, a process would be killed by the OOM killer
    - based on badness of each process
    - calculated based on the `vmsize` and uptime of each process

# To Small to Fail: Example

```
#define SMALL_SIZE 1024L
void exhaustheap() {
    for(long i=0; i < INT64_MAX; i++) {
        char* p = malloc (SMALL_SIZE);
        if(p == 0){
            printf("the %ldth malloc failed\n", i);
            break;
        } else {
            printf("access the %ldth memory chunk,...", i);
            memset (p, 0, sizeof (SMALL_SIZE));
            printf(", done\n", i);
        }
    }
}
```

```
#: sudo sysctl -w vm.overcommit_memory=2
#:~/4-memoxhaustion$ ./a.out
...
access the 2705176th memory chunk,..., done
the 2705177th malloc failed
```

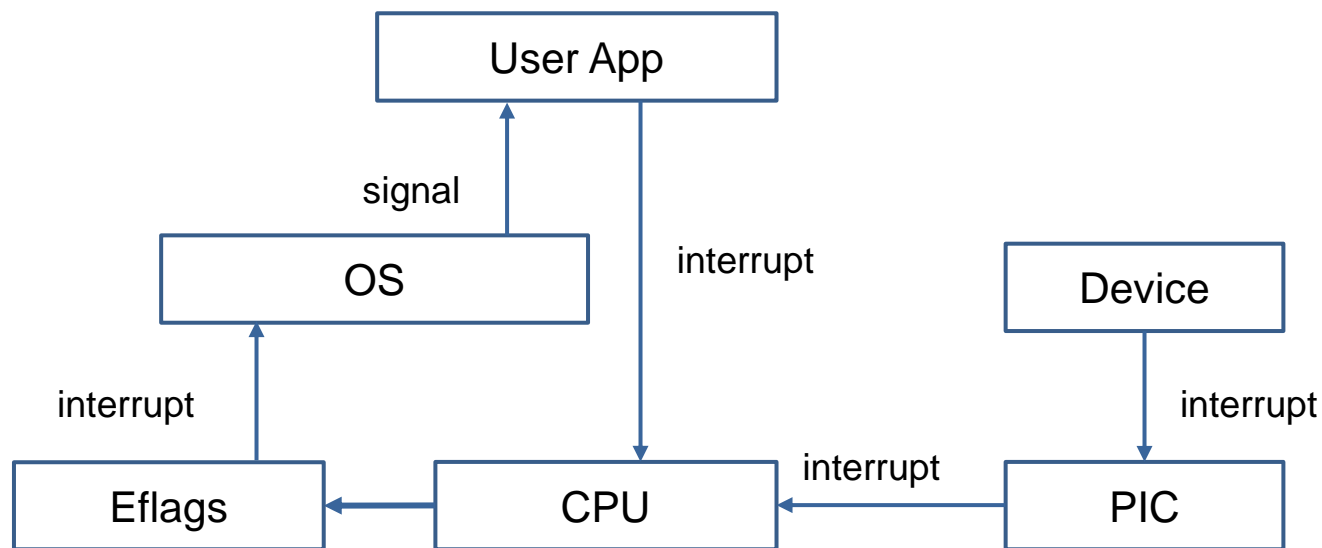
```
#: sudo sysctl -w vm.overcommit_memory=1
#:~/4-memoxhaustion$ ./a.out
...
access the 9013022th memory chunk,..., done
Killed
```

# 3. Exception Handling

---

# Exceptions based on Origin

- CPU: interrupt
- OS: signal
- Application: user-defined exceptions



# CPU Interrupt

- Page fault, divided by zero, etc
- Jump to the target exception handling address based on an interrupt vector, e.g., for X86
  - 0x00 Division by zero
  - 0x01 Single-step interrupt (see trap flag)
  - 0x03 Breakpoint (INT 3)
  - 0x04 Overflow
  - 0x06 Invalid Opcode
  - 0x0B Segment not present
  - 0x0C Stack Segment Fault
  - 0x0D General Protection Fault
  - 0x0E Page Fault
  - 0x10 x87 Floating Point Exception

# OS Signal

- Kernel sends to other processes (IPC)
- POSIX signals
  - SIGFPE: floating-point error, overflow, underflow...
  - SIGSEGV: segmentation fault, invalid address...
  - SIGBUS: bus error, memory alignment issue
  - SIGILL: illegal instruction
  - SIGABRT: abort
  - SIGKILL:
  - ...



# Register the OS Signal

- Register the OS signal with signal or sigaction

```
void sethandler(void (*handler)(int, siginfo_t *, void *)){
    struct sigaction sa;
    sa.sa_sigaction = handler;
    sigaction(SIGFPE, &sa, NULL);
}

void handler(int signo, siginfo_t *info, void *extra){
    printf("SIGFPE received!!!\n");
    exit(-1);
}

int main(void){
    sethandler(handler);
    int a = 0;
    int x = 100/a;
}
```

# Exception Handling Issue

- Where should the process continue?
  - find a landing pad
- How to set the required execution context?
  - restore callee-saved registers: rbp、rsp、rbx、r12-r15
- Release acquired resources
  - e.g, heap, file discriptor

# setjmp/longjmp



- `setjmp(env)`:
  - backup registers and sets a recover point
  - return 0 if called directly, otherwise return a value if called by `longjmp()`
- `longjmp(env,value)`:
  - jump to a target address determined by value
  - restore all callee-saved registers: `rbp`, `rsp`, `rbx`, `r12-r15`

```
static jmp_buf buf;
void handler(int signo, siginfo_t *info, void *extra){
    printf("SIGFPE received!!!\n");
    longjmp(buf,1);
}

int main(void){
    sethandler(handler);
    int a = 0;
    if (!setjmp(buf))
        int x = 100/a;
    else
        printf("Continue execution after a longjmp.\n");
}
```

# Discussion

- How to support multiple setjmps/longjmps?




## 4、Stack Unwinding

---

# Problem

- Callee-saved registers should be restored
- Setjmp/longjmp is inconvenient or inefficient if widely used
- Can we have a better solution?

```
0x401130: push    %rbp
0x401131: mov     %rsp,%rbp
0x401134: sub     $0x10,%rsp
0x401138: mov     %edi,-0x8(%rbp)
0x40113b: cmpl    $0x0,-0x8(%rbp)
0x40113f: jne     0x401151
0x401145: movl    $0x1,-0x4(%rbp)
0x40114c: jmpq    0x40116d
0x401151: mov     -0x8(%rbp),%eax
0x401154: mov     -0x8(%rbp),%ecx
0x401157: sub     $0x1,%ecx
0x40115a: mov     %ecx,%edi
0x40115c: mov     %eax,-0xc(%rbp)
0x40115f: callq   0x401130
0x401164: mov     -0xc(%rbp),%ecx
0x401167: imul    %eax,%ecx
0x40116a: mov     %ecx,-0x4(%rbp)
0x40116d: mov     -0x4(%rbp),%eax
0x401170: add     $0x10,%rsp
0x401174: pop     %rbp
0x401175: retq
```



# DWARF

- Calculate the information required for recovering from each instruction during compilation
- Such data format (DWARF) and mechanism is defined in the standard of ABI
- The program unwinds the call stack iteratively
- Different from the dynamic solution with setjmp.
  - more convenient, throw/try/catch is based on DWARF
  - more efficient

# How Does DWARF Work?

- To recover the context of the caller, we should know whether callee-saved registers have been changed
- Such callee-saved registers should be saved on the stack
- Record the address of each callee-saved register



# Example

- Calculate the canonical frame address or CFA
  - Find all instructions related to stack expansion/reduction
- Record the address of callee-saved registers related to CFA

```
push    %rbp
mov     %rsp,%rbp
sub     $0x10,%rsp
mov     %edi,-0x8(%rbp)
cmpl    $0x0,-0x8(%rbp)
jne     0x401151
movl    $0x1,-0x4(%rbp)
jmpq    0x40116d
mov     -0x8(%rbp),%eax
mov     -0x8(%rbp),%ecx
sub     $0x1,%ecx
mov     %ecx,%edi
mov     %eax,-0xc(%rbp)
callq   0x401130
mov     -0xc(%rbp),%ecx
imul    %eax,%ecx
mov     %ecx,-0x4(%rbp)
mov     -0x4(%rbp),%eax
add     $0x10,%rsp
pop     %rbp
retq
```

return address = CFA-8

CFA = cur rsp + 16, old rbp = CFA - 16,

CFA = cur rsp + 32

CFA = cur rsp + 16;

CFA = cur rsp - 8, old rbp is already restored



# Usage of DWARF

- Debugging: developers can obtain the call stack with `backtrace()`
- Exception handling: require further information to determine the landing pad or language specific information (personality routine)
  - C++ try-throw-catch
  - Rust stack unwinding

# In-Class Practice

- Handle the stack overflow issues of the following code
- Your task is to implement the sethandler and handler functions
  - Useful APIs: setjmp/longjmp, sigaction, sigaltstack
  - ref: <https://man7.org/linux/man-pages/man2/sigaltstack.2.html>

```
struct List{
    int val;
    struct List* next;
};
void process(struct List* list, int cnt){
    if(list->next != NULL)
        process(list->next, ++cnt);
}
void sethandler(void (*handler)(int, siginfo_t *, void *))
void handler(int signo, siginfo_t *info, void *extra);
void main(void){
    sethandler(handler);
    struct List* list = malloc(sizeof(struct List));
    list->val = 1;
    list->next = list;
    if (setjmp(buf) == 0)
        process(list, 0);
    else
        printf("Continue after segmentation fault\n");
}
```