

Lecture 5: Concurrency

XU, Hui

xuh@fudan.edu.cn



Outline

1. Risks of Concurrent Programs
2. Atomicity and Lock
3. Synchronization and Memory Barrier

1. Risks of Concurrent Programs

Risks of Concurrent Programs

- Data race or shared access
- Deadlocks
- Out-of-order execution
 - Compiler issue
 - CPU issue

An Example of Data Race

```
#define NUM 100
int global_cnt = 0;


void *mythread(void *in) {
    for (int i=0; i<NUM; i++)
        global_cnt++; //concurrently accessed by multiple threads
}

int main(int argc, char** argv) {
    pthread_t tid[NUM];
    for (int i=0; i<NUM; i++){
        assert(pthread_create(&tid[i], NULL, mythread, NULL)==0);
    }
    for (int i=0; i<NUM; i++){
        pthread_join(tid[i], NULL);
    }
    assert(global_cnt==NUM*NUM); //assertion could fail!
}
```

Hint of experiments: do not turn on optimization

Typical Scenario of Data Race

- Multiple threads access the same memory unit concurrently
- At least one access is nonatomic (write)
- For example, add operation is not atomic in X86 (CISC)
 - load-add-store (multiple instructions or micro ops)



X86 vs ARM/RISC-V

```
void toy(int x, int y){  
    int z = x+y;  
}
```



X86: operands can be mem

```
push    rbp  
mov     rbp, rsp  
mov     DWORD PTR [rbp-0x4],edi  
mov     DWORD PTR [rbp-0x8],esi  
mov     eax, DWORD PTR [rbp-0x4]  
add     eax, DWORD PTR [rbp-0x8]  
mov     DWORD PTR [rbp-0xc], eax  
pop     rbp  
ret
```

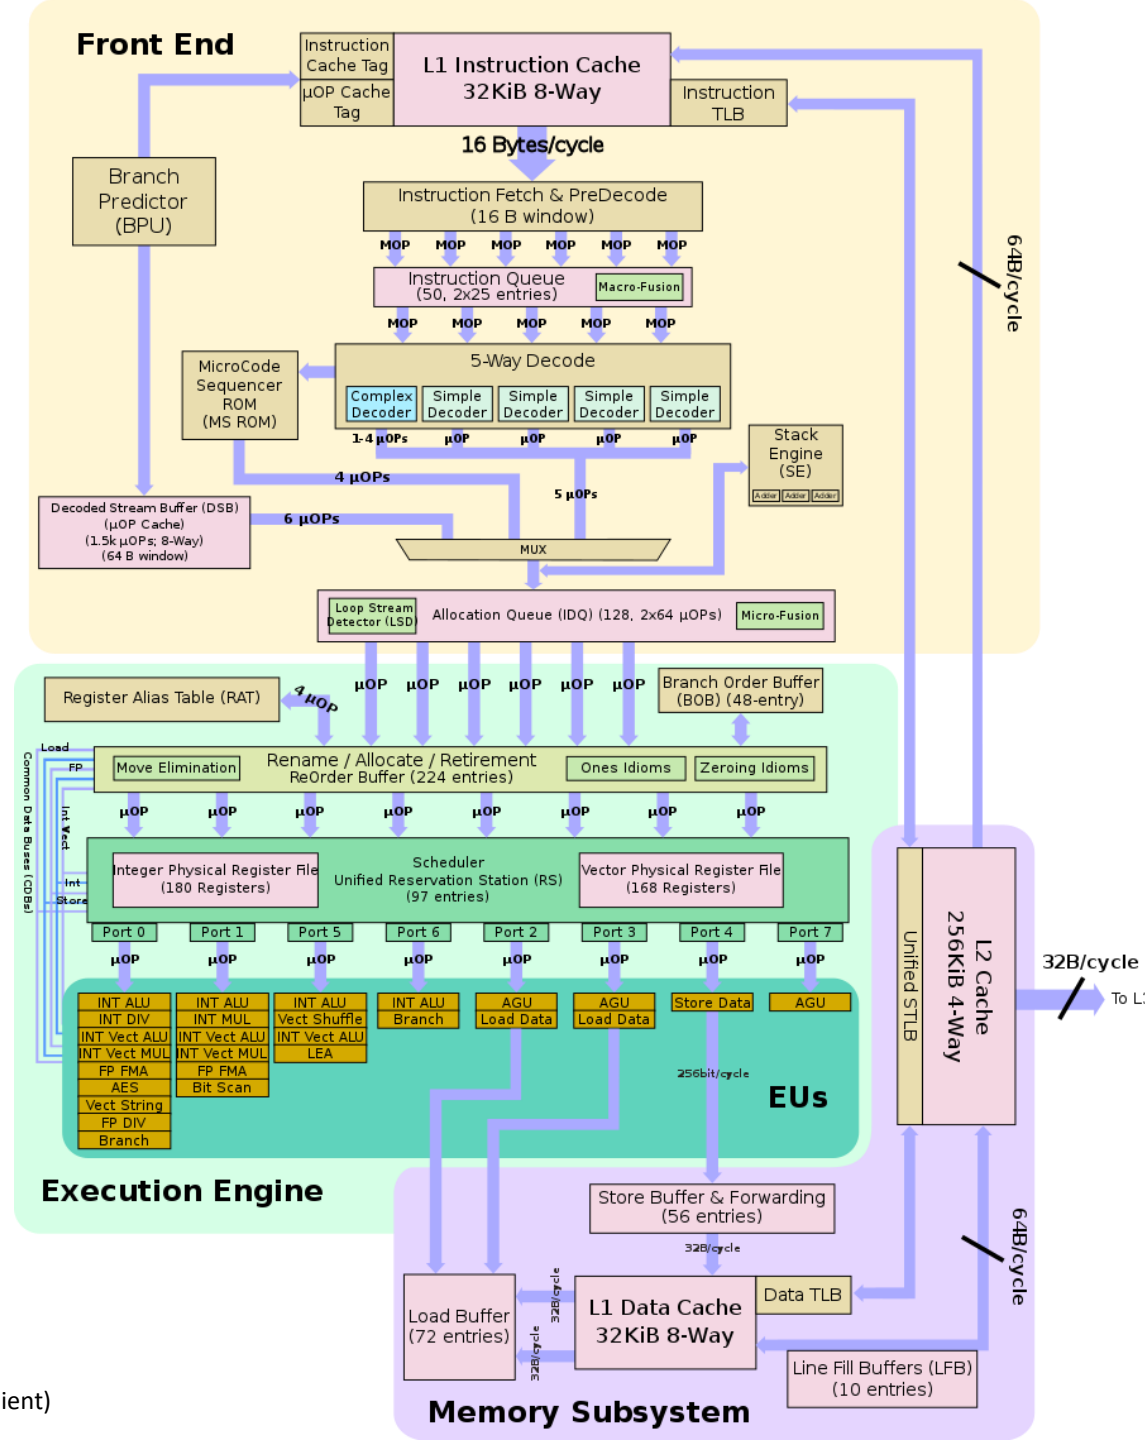


RISC-V: only registers can be
used as operands

```
addi    sp, sp, -32  
sd      ra, 24(sp)  
sd      s0, 16(sp)  
addi    s0, sp, 32  
sw      a0, -20(s0)  
sw      a1, -24(s0)  
lw      a0, -20(s0)  
lw      a1, -24(s0)  
addw    a0, a0, a1  
sw      a0, -28(s0)  
ld      ra, 24(sp)  
ld      s0, 16(sp)  
addi    sp, sp, 32  
ret
```

A close-up photograph of an Intel Core i7-6700K processor. The processor is a square, silver-colored integrated circuit mounted on a green printed circuit board (PCB). The top surface of the processor features the Intel logo, the text "INTEL® CORE™ i7", the model number "i7-6700K", the specification "SR2L0 4.00GHZ", and the cache size "X611A978" followed by a circled "64". The processor is surrounded by a dense array of gold-plated pins (BGA) used for mounting it on the motherboard. The green PCB has some white text and markings, including "000000" and "000000" at the top and bottom edges, and "000000" on the right edge.

- <https://en.wikichip.org/wiki/intel/microarchitectures/skylake> (client)



Question

- If the program runs on a single CPU core, will it still suffer race condition?

```
taskset -c 0 ./a.out
```

```
taskset -c 0 bash -c 'for i in {1..1000}; do ./a.out; done'
```

2. Atomicity and Lock

Atomic Instructions

- One instruction directly operates on the memory
 - Do not load the variable to the register
- Lock prefix guarantees atomicity of Micro Ops
 - X86 provides a “lock” prefix to achieve atomicity

```
mov    eax, DWORD PTR [rbp-0x14]
add    eax, 0x1
mov    DWORD PTR [rbp-0x14], eax
```



```
lock add DWORD PTR [rip+0x2ed3],0x1
```

Atomic Version


```
#define NUM 100
atomic_int global_cnt; //declaring the variable as atomic

void *mythread(void *from) {
    //use atomic API
    //__atomic_fetch_add(&global_cnt, 1, __ATOMIC_SEQ_CST);
    for (int i=0; i<NUM; i++)
        global_cnt++;
}

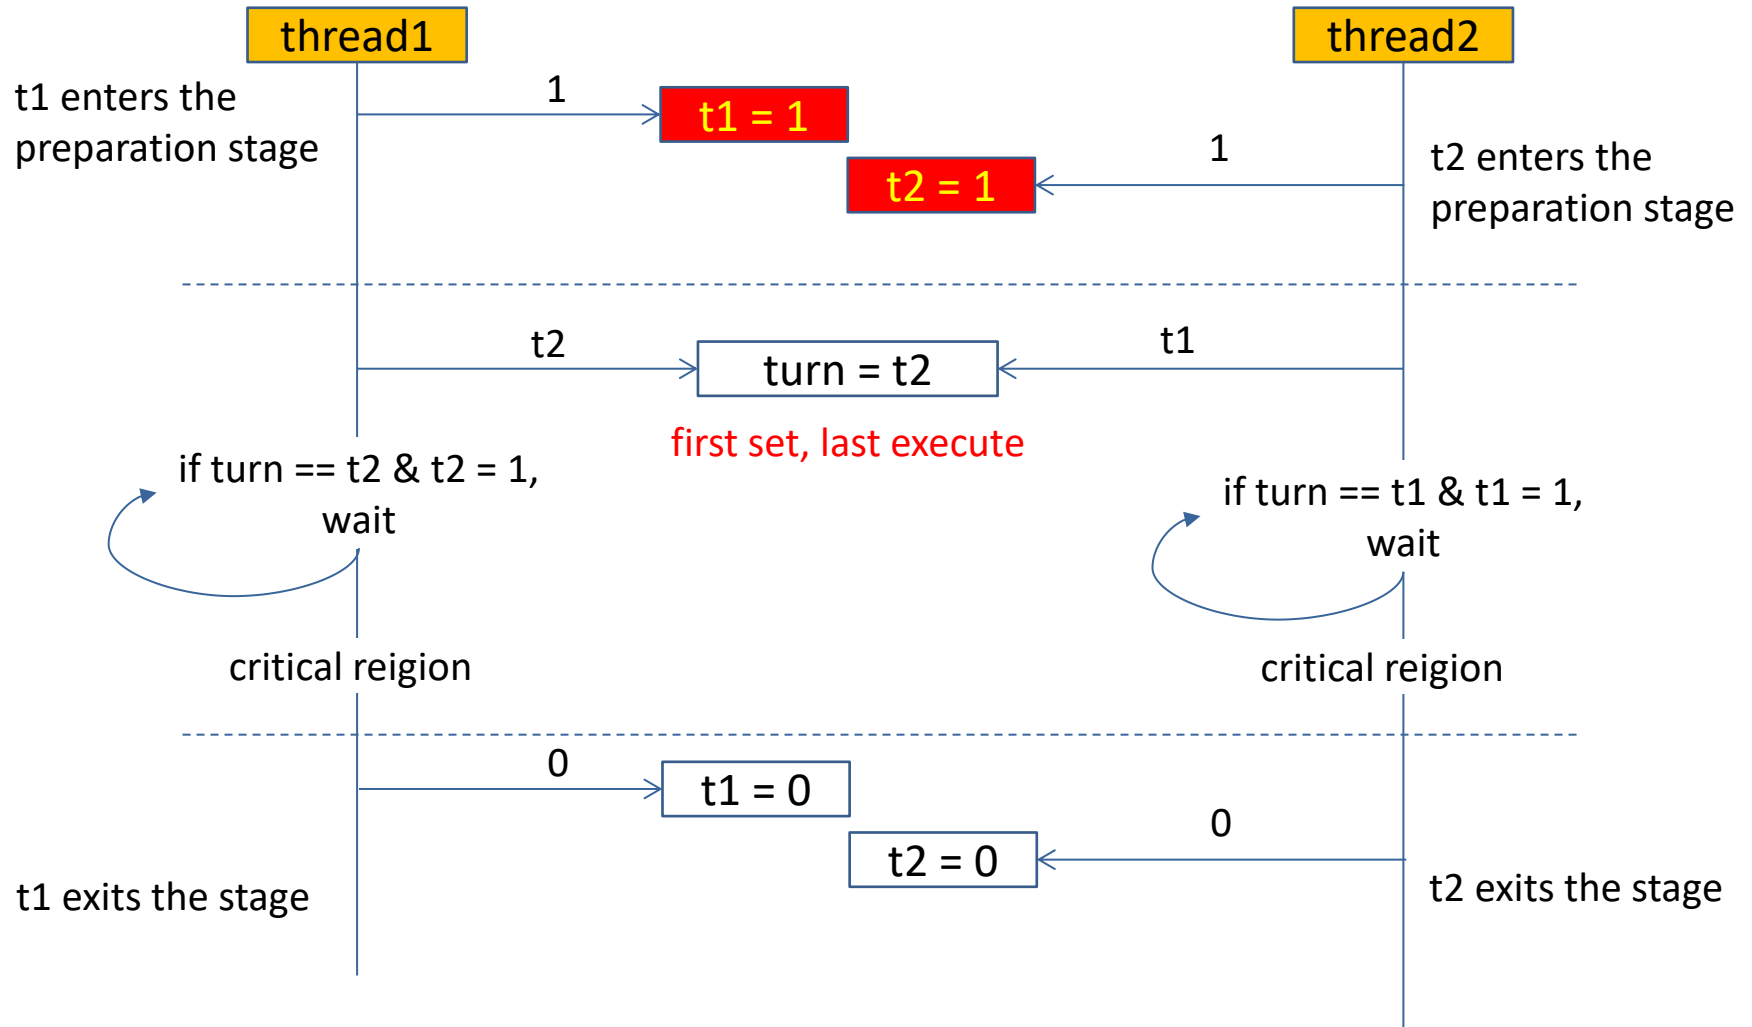
int main(int argc, char** argv) {
    pthread_t tid[NUM];
    for (int i=0; i<NUM; i++){
        assert(pthread_create(&tid[i], NULL, mythread, NULL)==0);
    }
    for (int i=0; i<NUM; i++){
        pthread_join(tid[i], NULL);
    }
    assert(global_cnt==NUM*NUM);
}
```

Mutual Exclusion or Mutex

- How to achieve atomicity for a sequence of code?
 - Entering the critical region without interference
- It is impossible with the "lock add" instruction.



Peterson Algorithm's for Mutex



Sample Code of Peterson's Algorithm

```
void* t0(void *from) {  
    flag[0] = true;  
    turn = 1;  
    while(flag[1]==true && turn==1)  
        sleep(1);  
    do_critical();  
    flag[0] = false;  
}
```

→ Give way to t1

```
void* t1(void *from) {  
    flag[1] = true;  
    turn = 0;  
    while(flag[0]==true && turn==0)  
        sleep(1);  
    do_critical();  
    flag[1] = false;  
}
```

→ Give way to t0

Mutex base on Atomic Instructions

- How to achieve atomic compare and set/swap?
 - x86 instruction: cmpxchg
 - C API: atomic_compare_exchange_strong

```
# based on rax  
lock cmpxchg dst src
```

sematic



```
eax = 0, src = 1
```

```
if(dst == eax) {  
    dst = src;  
}  
else {  
    eax = dst;  
}
```

- dst is the lock flag
- rax is the return value of check

```
atomic_compare_exchange_strong(&dst, &test, src)
```

exactly the same with cmpxchg

Type of Locks: Mutex Lock vs Read-Write Lock

- Mutex lock: only one thread can access a shared resource at a time.
 - Lock Acquisition (lock) – A thread requests ownership of the mutex.
 - Critical Section – The thread executes its operations on the shared resource.
 - Lock Release (unlock) – The thread releases the mutex, allowing others to proceed.
- Read-write lock: multiple threads to read concurrently; only one thread to write exclusively.
 - Unlocked – No threads hold the lock.
 - Read(n) – Multiple threads can read at the same time (n readers).
 - Write (exclusive) – Only one thread can write, and no readers are allowed.

Type of Locks: Spin Lock

- Thread tries to acquire the lock
- If the lock is free, it proceeds
- If the lock is held, the thread keeps checking (spinning) until it's released

Type of Locks: Pessimistic vs Optimistic Lock

- Pessimistic Lock: Assumes that conflicts will occur, so it locks the resource before accessing it, e.g., mutex lock.
 - A thread acquires a lock before accessing shared data.
 - Other threads must wait until the lock is released.
- Optimistic Lock: Assumes conflicts are rare, so it accesses the resource without locking.
 - Read the data.
 - Do some work.
 - Check if the data has changed (validation).
 - If unchanged, commit the update.
 - If changed by another thread, retry or abort.

Question

- Is `shared_ptr` of C++ thread-safe?
 - reference counter
 - data read/write

3. Synchronization and Memory Barrier

Out-of-Order Execution

- Compiler reordering during optimization
- CPU out-of-order execution
- This lecture focuses on compile-time reordering

Compiler Reordering

- Supposing optimization (*e.g.*, -O2 or O3) is enabled, the compiler might make mistakes.

```
int a = 1;
```

```
void *t0 (void* in){  
    while (a);  
}
```

```
void *t1 (void* in){  
    a = 0;  
}
```

```
0x00401150 <+0>:    cmp     DWORD PTR [rip+0x2ee9], 0x0  
0x00401157 <+7>:    je      0x401162 <t0+18>  
0x00401159 <+9>:    nop     DWORD PTR [rax+0x0]  
0x00401160 <+16>:    jmp     0x401160 <t0+16>  
0x00401162 <+18>:    ret
```



infinite loop

Another Example

- The following assertion could fail on some platforms if the execution order cannot be guaranteed.

```
atomic_int a = 1;  
atomic_int b = 1;
```

```
void *t0 (void* in){  
    a = 0;  
    b = 0;  
}
```

```
void *t1 (void* in){  
    while(!b);  
    assert(!a);  
}
```


Use Memory Barrier (Fence)

An instruction on x86 processors

Arguments and ret value

```
#define barrier() __asm__ __volatile__("mfence": : : "memory");
```

Tells the compiler that this instruction affects memory, preventing reordering of memory operations across this barrier

Example

- Discard all variable values on registers
 - Reload them from memory
- Guarantee happens-before: operations prior to the barrier are always executed before operations after the barrier.

```
void *t0 (void* in){  
    while (a)  
        barrier();  
}
```

```
void *t0 (void* in){  
    a = 0;  
    barrier();  
    b = 0;  
}
```

Relax the Synchronization Requirement

- We only want some variables to be updated.
- Use volatile when declaring a variable.
- Do not prevent reordering.

```
volatile int a = 1;  
void *t0 (void* in){  
    while (a) ;  
}
```

Further Relax the Requirement

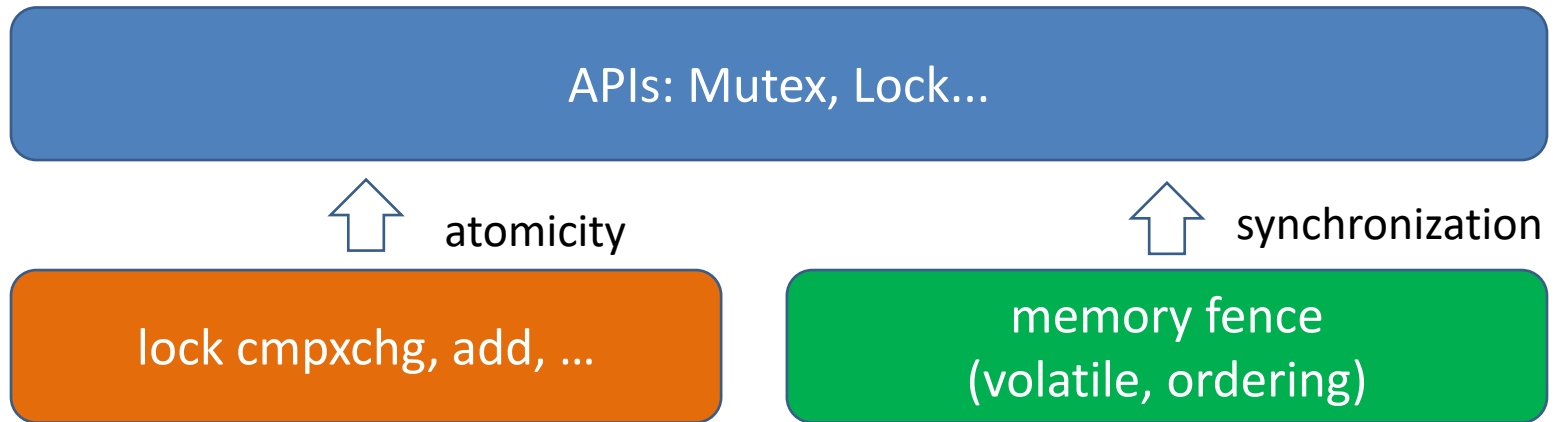
- We want the value to be updated in specific program points.
 - Use ACCESS_ONCE(), which is also based on volatile.
- Further relax the restrictions based on operations:
 - READ_ONCE and WRITE_ONCE()

```
#define ACCESS_ONCE(x) (*(volatile typeof(x) *)&(x))
volatile int a = 1;
void *t0 (void* in){
    while (ACCESS_ONCE(a))
        ;
}
```

Relax the Happens-Before Requirement

- Use specific memory ordering
 - Sequential consistency (default on x86)
 - The most strong one, no reordering across the barrier.
 - Acquire-release: commonly used for locks
 - Acquire: no reads or writes can be reordered before this load.
 - Release: no reads or writes can be reordered after this store.
 - Relaxed
 - No synchronization or ordering constraints, only atomicity.

Summary



In-Class Practice

- 1) Demonstrate that `std::shared_ptr` in C++ is not thread-safe.
 - Hint: Design code that triggers an error by causing a race condition in the reference counter update.
- 2) (Optional) Modify the shared pointer you implemented last week to make it thread-safe.