

Lecture 6: Concurrency

XU, Hui

xuh@fudan.edu.cn



Outline

1. Risks of Concurrent Programs
2. Atomicity and Lock
3. Synchronization and Memory Barrier

1. Risks of Concurrent Programs

Risks of Concurrent Programs

- Data race or shared access
- Deadlocks
- Out-of-order execution
 - Compiler issue
 - CPU issue

An Example of Data Race

```
#define NUM 100
int global_cnt = 0;


void *mythread(void *in) {
    for (int i=0; i<NUM; i++)
        global_cnt++; //concurrently accessed by multiple threads
}

int main(int argc, char** argv) {
    pthread_t tid[NUM];
    for (int i=0; i<NUM; i++){
        assert(pthread_create(&tid[i], NULL, mythread, NULL)==0);
    }
    for (int i=0; i<NUM; i++){
        pthread_join(tid[i], NULL);
    }
    assert(global_cnt==NUM*NUM); //assertion could fail!
}
```

Hint of experiments: do not turn on optimization

Typical Scenario of Data Race

- Multiple threads access the same memory unit concurrently
- At least one access is nonatomic (write)
- For example, add operation is not atomic in X86 (CISC)
 - load-add-store (multiple instructions or micro ops)



X86 vs ARM/RISC-V

```
void toy(int x, int y){  
    int z = x+y;  
}
```



X86: operand should be mem

```
push    rbp  
mov     rbp, rsp  
mov     DWORD PTR [rbp-0x4],edi  
mov     DWORD PTR [rbp-0x8],esi  
mov     eax, DWORD PTR [rbp-0x4]  
add     eax, DWORD PTR [rbp-0x8]  
mov     DWORD PTR [rbp-0xc], eax  
pop     rbp  
ret
```

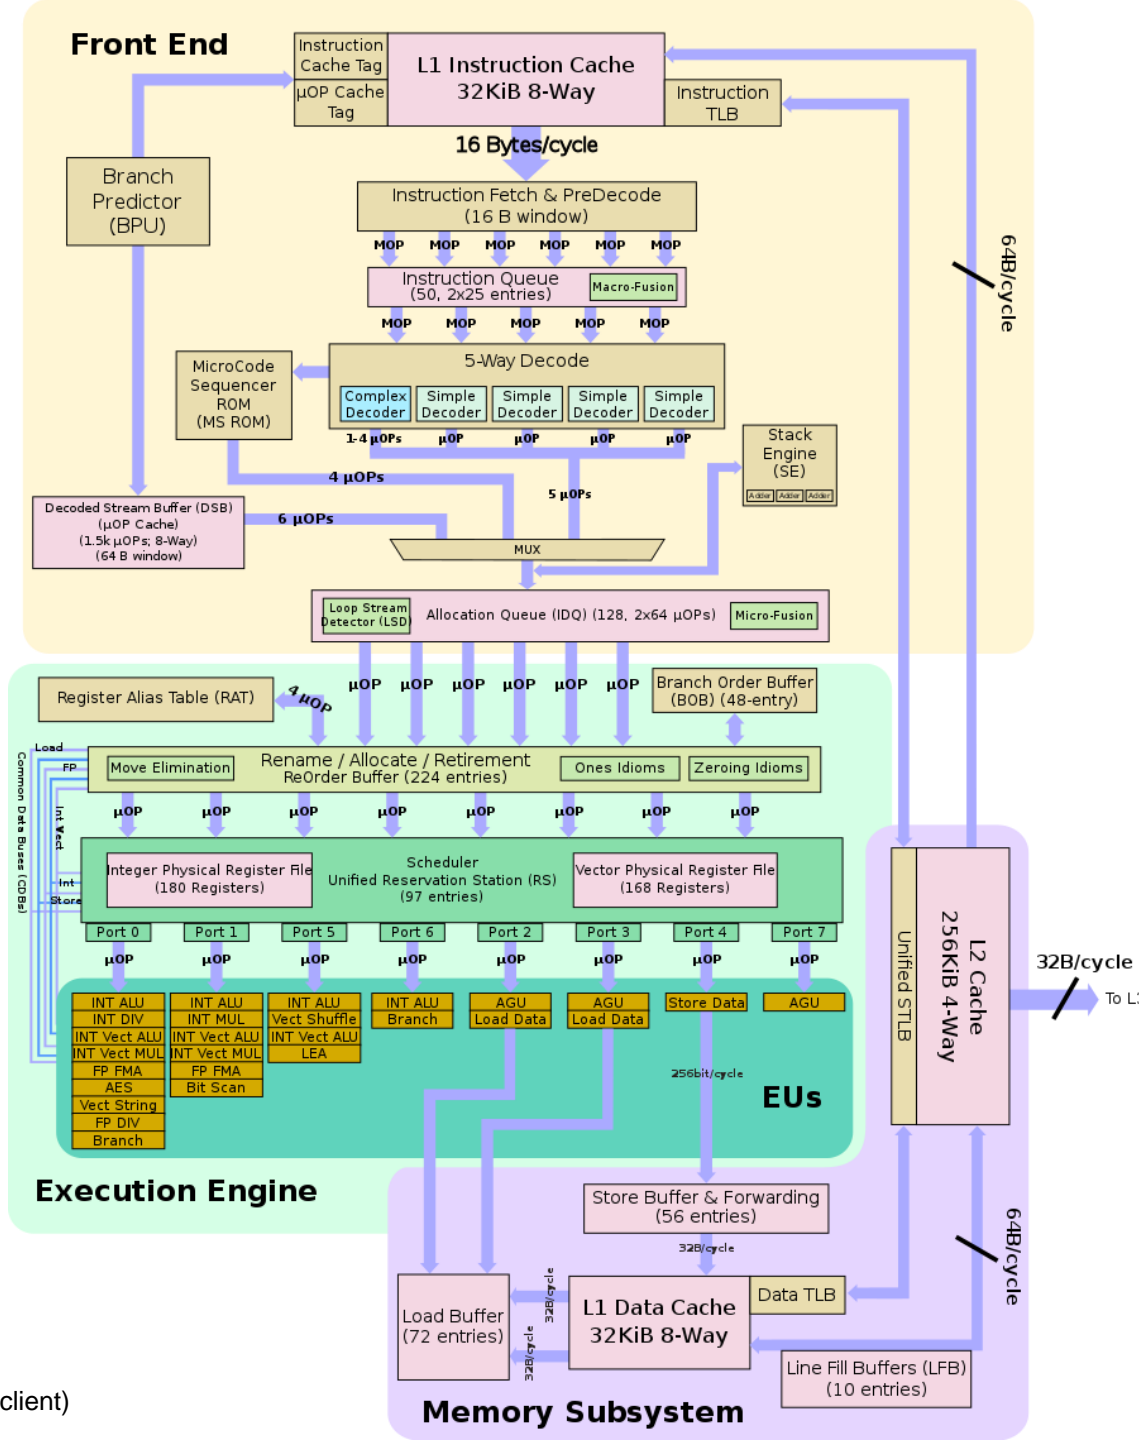


RISC-V: only registers can be used as operands

```
addi    sp, sp, -32  
sd      ra, 24(sp)  
sd      s0, 16(sp)  
addi    s0, sp, 32  
sw      a0, -20(s0)  
sw      a1, -24(s0)  
lw      a0, -20(s0)  
lw      a1, -24(s0)  
addw    a0, a0, a1  
sw      a0, -28(s0)  
ld      ra, 24(sp)  
ld      s0, 16(sp)  
addi    sp, sp, 32  
ret
```

A photograph of an Intel Core i7-6700K processor. The processor is a square, silver-colored integrated circuit mounted on a green printed circuit board (PCB). The top surface of the processor is marked with the Intel logo, "INTEL® CORE™ i7", "i7-6700K", "SR2L0 4.00GHZ", and "X611A978" followed by the Intel logo. The processor is connected to the PCB via numerous gold-plated pins.

- [https://en.wikichip.org/wiki/intel/microarchitectures/skylake_\(client\)](https://en.wikichip.org/wiki/intel/microarchitectures/skylake_(client))



2. Atomicity and Lock

Use Atomic Instructions

- One instruction directly operates on the memory
 - do not load the variable to the register
- Lock prefix guarantees atomicity of Micro Ops
 - X86 provides a “lock” prefix for atomicity

```
gef> disass mythread
```

```
Dump of assembler code for function mythread:
```

```
...
```

```
0x0040117a <+42>:    mov     eax, DWORD PTR [rbp-0x14]
```

```
0x0040117d <+45>:    add     eax, 0x1
```

```
0x00401180 <+48>:    mov     DWORD PTR [rbp-0x14],eax
```

```
...
```



```
lock add DWORD PTR [rip+0x2ed3],0x1
```

Atomic Version


```
#define NUM 100
atomic_int global_cnt; //way 1: by declaring the variable as atomic

void *mythread(void *from) {
    //way 2: use atomic API
    //__atomic_fetch_add(&global_cnt, 1, __ATOMIC_SEQ_CST);
    for (int i=0; i<NUM; i++)
        global_cnt++;
}

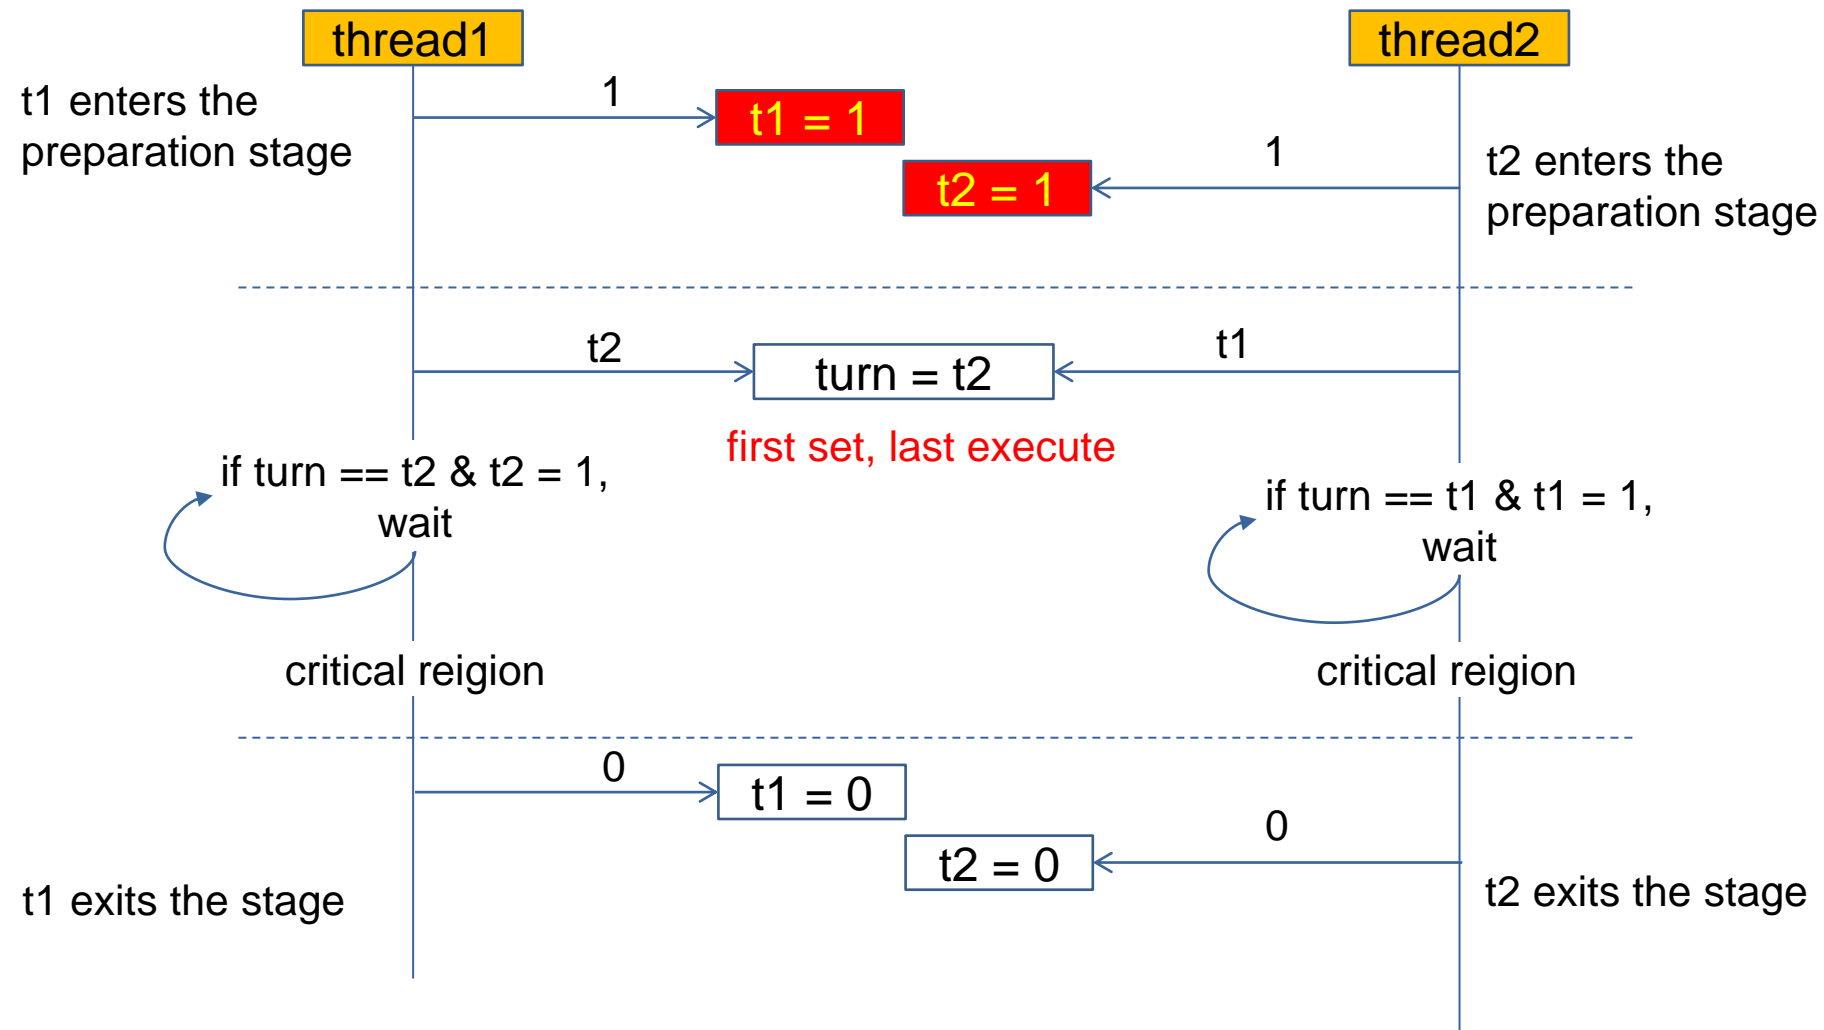
int main(int argc, char** argv) {
    pthread_t tid[NUM];
    for (int i=0; i<NUM; i++){
        assert(pthread_create(&tid[i], NULL, mythread, NULL)==0);
    }
    for (int i=0; i<NUM; i++){
        pthread_join(tid[i], NULL);
    }
    assert(global_cnt==NUM*NUM);
}
```

Mutual Exclusion or Mutex

- How to achieve atomicity for a sequence of code?
 - entering the critical region without interference
- It is impossible with the "lock add" instruction



Peterson Algorithm's for Mutex



Sample Code of Peterson's Algorithm

```
void* t0(void *from) {
    flag[0] = true;
    turn = 1;
    while(flag[1]==true && turn==1)
        sleep(1);
    do_critical();
    flag[0] = false;
}
```

```
void* t1(void *from) {
    flag[1] = true;
    turn = 0;
    while(flag[0]==true && turn==0)
        sleep(1);
    do_critical();
    flag[1] = false;
}
```

```
int flag[2], turn;
int x=0;

void do_critical(){
    x++;
}

int main(int argc, char** argv) {
    pthread_t tid[2];
    assert(pthread_create(&tid[0], NULL, t0, NULL)==0);
    assert(pthread_create(&tid[1], NULL, t1, NULL)==0);
    pthread_join(tid[0], NULL);
    pthread_join(tid[1], NULL);
}
```

Use Atomic Instructions

- How to achieve atomic compare and set/swap?
 - x86 instruction: `cmpxchg`
 - C API: `atomic_compare_exchange_strong`

based on rax
`lock cmpxchg dst src`

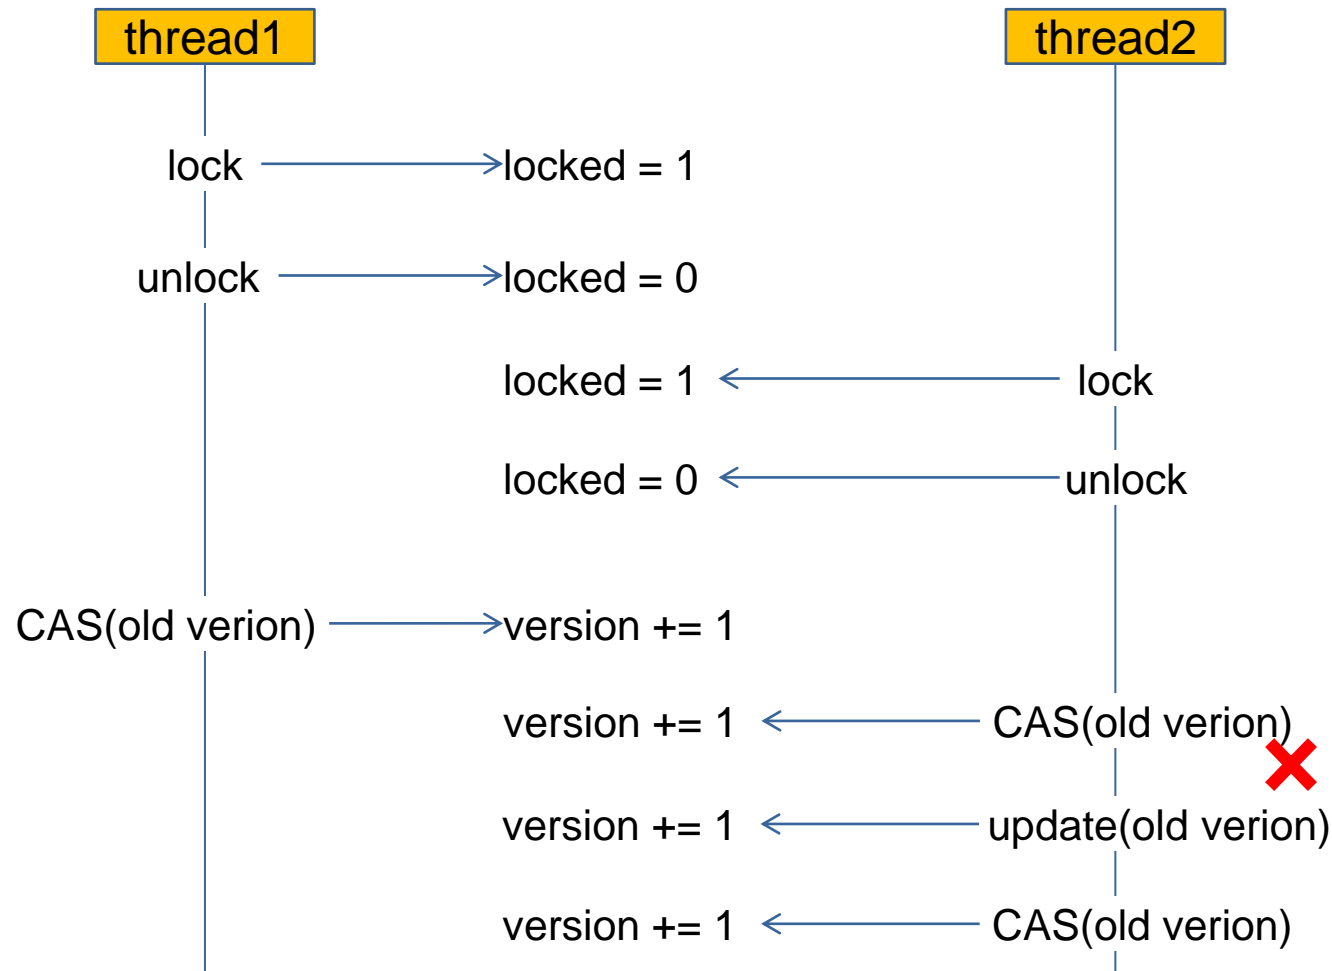
Semantic:

```
if(dst == eax) {  
    dst = src;  
    ZERO_FLAG = 1;  
}  
else {  
    eax = dst;  
    ZERO_FLAG = 0;  
}
```

`atomic_compare_exchange_strong(&dst, &test, src)`

exactly the same with `cmpxchg`

Pessimistic Lock vs Optimistic Lock



More Types of Locks

- Mutex lock
- Spin lock: continuously try to acquire the lock
- Read-write lock:
 - State: write/read(n)/unlocked
 - A variable can be read by multiple thread concurrently.
 - Weakness: cannot guarantee serializability of operations

Question

- Is `shared_ptr` of C++ thread-safe?
 - reference counter
 - data read/write
- Can you implement a thread-safe construct?
- We will learn `ARC<Mutex<T>>` in Rust

3. Synchronization and Memory Barrier

Out-of-Order Execution

- Compiler reordering during optimization
- CPU out-of-order execution
- This lecture focuses on compile-time ordering

Compiler Reordering

- Supposing optimization (e.g., -O2 or O3) is enabled, the compiler might make mistakes.

```
atomic_int a = 1;
```

```
void *t0 (void* in){  
    while (a);  
}
```

```
void *t1 (void* in){  
    a = 0;  
}
```

```
0x00401150 <+0>:      cmp     DWORD PTR [rip+0x2ee9], 0x0  
0x00401157 <+7>:      je      0x401162 <t0+18>  
0x00401159 <+9>:      nop     DWORD PTR [rax+0x0]  
0x00401160 <+16>:  jmp     0x401160 <t0+16>  
0x00401162 <+18>:  ret
```



infinite loop

Another Example

- The following assertion could fail on some platforms if the execution order cannot be guaranteed

```
atomic_int a = 1;  
atomic_int b = 1;
```

```
void *t0 (void* in){  
    a = 0;  
    b = 0;  
}
```

```
void *t1 (void* in){  
    while(!b);  
    assert(!a);  
}
```

Use Memory Barrier (Fence)

- Discard all variable values on registers
 - Reload them from memory
- Guarantee happens-before: operations prior to the barrier are always executed before operations after the barrier.

```
#define barrier() __asm__ __volatile__("" : : : "memory");
```

```
void *t0 (void* in){  
    while (a)  
        barrier();  
}
```

```
void *t0 (void* in){  
    a = 0;  
    barrier();  
    b = 0;  
}
```

Relax the Variables

- We only want some variables to be updated:
 - use volatile when declaring a variable
- We only want the variable to be updated in specific program points:
 - Use ACCESS_ONCE(), which is also based on volatile
- Further relax the restrictions based on operations:
 - READ_ONCE and WRITE_ONCE()

```
volatile int a = 1;
void *t0 (void* in){
    while (a) ;
}
```

```
#define ACCESS_ONCE(x) (*(volatile typeof(x) *)&(x))
volatile int a = 1;
void *t0 (void* in){
    while (ACCESS_ONCE(a)) ;
}
```


Relax Happens-Before Requirement

- Use specific memory ordering
 - Sequential consistency (default on x86)
 - the most strong one, no reordering across the barrier;
 - Acquire-release
 - release: no reads or writes in the current thread can be reordered after this store
 - acquire: no reads or writes in the current thread can be reordered before this load
 - commonly used for locks
 - Relaxed
 - no synchronization or ordering constraints
 - only atomicity

Summary



In-Class Practice

- 1) Implement locks in C and demonstrate the effectiveness
 - a mutex (lock)
 - an optimistic lock
 - based on the `atomic_compare_exchange_weak()` API, https://en.cppreference.com/w/c/atomic/atomic_compare_exchange
- 2) Optional: try to implement a thread-safe shared pointer for a data structure with mutex member functions.