

Lecture 6: Rust Ownership

Hui Xu

xuh@fudan.edu.cn



Outline

- 1. Ownership
- 2. RAI and Lifetime
- 3. Unsafe Code

1. Ownership

Motivation of Design

- Dangling pointer is unacceptable.
- Causal of dangling pointer?
 - False reclaim: manual heap management with malloc/free
 - Automatic reclaim?
 - Static analysis is impossible because alias analysis is NP-hard
 - Shared pointer or garbage collection is inefficient
- Rust comes to rescue.
 - Static analysis without NP-hard problems.
 - How? Ownership-based heap management.

Overall Idea of Ownership

- Each object is owned by one variable.
 - Recall C++ unique pointer
- Ownership can be borrowed in two modes:
 - immutable: read only
 - mutable: read/write
- Exclusive mutability principle: two variables should not share mutable access to the same object at one program point.

Ownership & Borrowing

- Borrowed ownership will be returned automatically if no longer used.

```
let mut alice = Box::new(1);
```

→ alice owns the object

```
let bob = alice;
```

→ alice transfers the ownership to bob;

```
println!("bob:{}", bob);
```

```
println!("alice:{}", alice);
```



```
let mut alice = Box::new(1);
```

```
let bob = &alice;
```

→ bob borrows the ownership

```
println!("bob:{}", bob);
```

→ bob returns the ownership to
alice automatically

```
println!("alice:{}", alice);
```



Violating Exclusive Mutability

```
let mut alice = 1;  
let bob = &mut alice;  
println!("bob:{}", bob);  
println!("alice:{}", alice);
```

mutable borrow

bob returns the ownership

```
let mut alice = 1;  
let bob = &mut alice;  
println!("alice:{}", alice);  
println!("bob:{}", bob);
```

violate exclusive mutability

Move Operator (=)

- If a type is not Copy (trait), move transfers the ownership.
 - e.g., Box<T> is not copy
- If a type is Copy, move does not transfer the ownership and only copies the value.

```
let mut alice = 1;  
let bob = alice;  
println!("bob:{}", bob);  
println!("alice:{}", alice);
```

→ alice owns the object

→ copy: duplicate the object



Which Types Can be Copy?

- Primitive types on stack.
- Compound types with all fields implementing Copy.
- How to (deep) copy objects of other non-Copy types?
 - By implementing the Clone trait.
 - Each Copy type is also Clone.

```
let mut alice = Box::new(1);
```

→ alice owns the object

```
let bob = alice.clone();
```

→ clone the object for bob

```
println!("bob:{}", bob);
```


```
println!("alice:{}", alice);
```




Mutability

mutable object

```
let mut alice = 1;  
alice+=1;
```




```
let alice = 1;  
alice+=1;
```




mutable borrow


```
let mut alice = 1;  
let bob = &mut alice;  
*bob+=1;
```



```
let mut alice = 1;  
let bob = &alice;  
*bob+=1;
```



```
let alice = 1;  
let bob = &mut alice;  
*bob+=1;
```



Mutability cont'd

```
let mut alice = 1;  
let mut carol = 1;  
let mut bob = &mut alice;  
*bob+=1;  
bob = &mut carol;  
*bob+=1;
```



→ mutable object bob + mutable borrow

```
let mut alice = 1;  
let mut carol = 1;  
let bob = &mut alice;  
*bob+=1;  
bob = &mut carol;  
*bob+=1;
```



→ immutable object bob + mutable borrow

→ bob cannot be modified

Why Ownership is Effective for Static Analysis

- Restrict the complexity of the alias analysis problem.
 - There is only one mutable pointer at each program point.
 - Only mutable pointers can lead to dangling pointers.
 - We only need to trace the mutable pointer for each object.
 - e.g., via a stack-based approach.

Pros and Cons

- Benefit: compile-time method, no runtime cost
- Cons: when shared mutability is a Must...
 - Such as double linked lists?
 - Alternatives: shared pointer or unsafe code.

2. RAI and Lifetime

RAI: Resource Acquisition is Initialization

Idea of RAI

- Tie resources to lifetime.
- Each object is allocated and initialized once created.
 - all pointers refer to particular objects
 - no raw or dangling pointers
 - no uninitialized memory
- Deallocation automatically when the object is dead.
 - no manual deallocation is needed
 - achieved through static lifetime inference

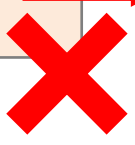
Lifetime

- Each object has a lifetime constraint.
- The object is reclaimed automatically after death.
- A variable cannot borrow an object with a shorter lifetime.

```
let alice;  
{  
    let bob = 5;  
    alice = &bob;  
}  
println!("alice:{}", alice);
```

→ bob lives in this subscope

→ alice points to an expired object



Review Move for Non-Copy Types

- Expand the lifetime of the object on heap

```
let alice;  
{  
    let bob = Box::new(1);  
    alice = bob;  
}  
println!("alice:{}", alice);
```

transfer the ownership to alice



```
fn testret() -> Box<u64>{  
    Box::new(1)  
}
```

move the ownership to the ret value

```
let r = testret();  
println!("return:{}", r);
```



Lifetime Declaration

- Lifetime should be declared during function declaration

```
fn longer<'a>(x:&'a String, y:&'a String)->&'a String{
    if x.len()>y.len(){
        x
    } else {
        y
    }
}
fn stringcmp(){
    let str1 = String::from("alice");
    let str2 = String::from("bob111");
    let result = longer(&str1, &str2);
    println!("The longer string is {}", result);
}
```

Partial Order of Lifetime

- `<'a: 'b, 'b>` means a is relatively larger than b

```
fn longer<'a: 'b, 'b>(x:&'a String, y:&'b String)->&'b String{  
    if x.len()>y.len(){  
        x  
    } else {  
        y  
    }  
}  
  
fn stringcmp(){  
    let str1 = String::from("alice");  
    let result;  
    //{  
        let str2 = String::from("bob111");  
        result = longer(&str1, &str2);  
    //}  
    println!("The longer string is {}", result);  
}
```

The return value cannot be 'a. Why?

Non-lexical Lifetime

- The default mode is non-lexical unless necessary.
- Rust compiler tries to minimize the lifespan.

```
'a: { let str1 = "alice";  
    'b: { let str2 = "bob";  
        'c: { let result = longer(str1,str2);  
              println!("The longer string is {}", result);  
        }  
    }  
}
```

Lifetime Elision to Be More Ergonomic

- Lifetime declaration can be elided if
 - there is only one input lifetime position
 - there are multiple positions, but one is `&self` or `&mut self`
 - assign the lifetime of `self` to elided output lifetimes

```
fn foo<'a>(s: &'a str, until: usize) -> &'a str;
```



```
fn foo(s: &str, until: usize) -> &str;
```



```
fn foo(s: &str, t: &str) -> &str; // illegal
```



```
fn foo<'a, 'b>(&'a mut self, t: &'b str) -> &str;
```



```
fn foo(&mut self, t: &str) -> &str;
```



More About Lifetime

- A static object means it lives for the entire lifetime.
 - use the reserved lifetime 'static.
 - all strings are static by default.
- The lifetime can be unbounded.
 - Similar to static but more flexible during type check

```
let s: &str = "hello world";
```



equivalent to

```
let s: &'static str = "hello world";
```


```
fn foo<'a>(s: *const String) -> &'a str;
```

unbounded lifetime

Automatic Reclaim/Drop

- Objects of Copy type (stack) can be reclaimed automatically.
- For other objects with heap data?
 - Drop (trait) unused objects by calling the destructor.
 - Recursively call the destructor of each field.
- Drop and Copy are exclusive in Rust.
 - Box<T> has the Drop trait, but not Copy

```
struct MyType {a:u8, b:Box<u64>}

impl Drop for MyType {  automatically executed when the lifetime ends
    fn drop(&mut self){
        println!("dropping MyType object...");
    }
}

fn testdrop(){
    let v = MyType {a:1, b:Box::new(2)};
}
```

Option<T> for Uninitialized Objects

- Option: an enumerate type
 - Some(T): the object type
 - None: if the object is uninitialized (null pointer)

```
pub enum Option<T> {  
    None,  
    Some(T),  
}  
  
let v = Some(...)  
match v.next {  
    Some(n) => ...,  
    None => panic!(),  
}
```


Example with a Singly-linked List

```
struct List {
    val: u64,
    next: Option<Box<List>>,
}
impl List {
    fn new(val: u64) -> Self {
        List { val, next: None }
    }
    fn prepend(self, val: u64) -> Self {
        List {
            val,
            next: Some(Box::new(self)),
        }
    }
    fn append(&mut self, val: u64) {
        let mut current = self;
        while let Some(ref mut next) = current.next {
            current = next;
        }
        current.next = Some(Box::new(List::new(val)));
    }
}
```

RAII for Thread Panic

- In a multi-threaded application, what happens when one thread exit exceptionally ?
 - abort: directly terminate the thread
 - panic: perform stack unwinding before exit
- Importance of RAII during stack unwinding
 - release locks (mutex)
 - release opened file descriptors
 - release allocated memories on heap

Sample Multi-thread Program


- When a spawned thread panics, unwind its stack.
- The main thread continues execution.
- Ineffective for fatal errors: *e.g.*, stack overflow

```
let handle = thread::spawn(|| {  
    for i in 0..5 {  
        println!("new thread print {}", i);  
        thread::sleep(Duration::from_millis(10));  
    }  
    panic!();  
    //recursive();  
});  
for i in 0..10 {  
    println!("main thread print {}", i);  
    thread::sleep(Duration::from_millis(10));  
}  
handle.join();
```

3. Unsafe Rust

Problem for Double-Linked List

- A node is owned by both its prev and next node.
- Violate exclusive mutability.



```
struct Node {  
    val: u64,  
    prev: Option<Box<Node>>,  
    next: Option<Box<Node>>,  
}
```

Unsafe Code

- Operations that may lead to undefined behaviors are unsafe.
 - dereference raw pointers
 - call unsafe functions
 - call functions of foreign language (FFI)
- Unsafe code can only be used with the unsafe marker.

```
let mut num = 5;  
let r1 = &num as *const i32;  
println!("r1 is: {}", unsafe { *r1 });
```

→ r1 is a raw pointer

→ raw pointer dereference

Dereference raw pointers

```
unsafe fn foo() {  
    let addr = 0x012345usize;  
    let r = addr as *const i32;  
}  
unsafe { foo(); }
```


→ define an unsafe function

→ call the unsafe function

Call unsafe functions


Trust Model

- Rust does not trust developers, so only safe code is allowed.
- If a developer declares he knows the risk, Rust will trust him.



Principle of Rust

- Safe APIs should not incur undefined behaviors.
- Interior unsafe: wrap unsafe code into safe APIs.
- Avoid using unsafe code unless necessary.



Unsafe Version with Raw Pointers

- The objects pointed by raw pointers are not owned.
- The resource may not be dropped automatically.
- It is also prone to dangling pointers (not RAII).

```
struct Node {  
    val: u64,  
    next: *mut List,  
    prev: *mut List,  
}  
  
impl Node {  
    fn new(val: u64) -> *mut Node {  
        Box::into_raw(Box::new(Node { val,  
            next: ptr::null_mut(), prev: ptr::null_mut() })))  
    }  
  
    unsafe fn prepend(head: *mut List, val: u64) -> *mut List {  
        let new_node = List::new(val);  
        if !head.is_null() {  
            (*new_node).prev = head;  
            (*head).next = new_node;  
        }  
        new_node  
    }  
}
```


Solution Hint with RC<T> and Weak<T>

- RC<T>/Weak<T>: single-thread reference-counting pointer
 - enable shared mutable aliases
- RefCell<T>: enable mutable access

```
struct Node {  
    val: u64,  
    prev: Option<Weak<RefCell<Node>>>,  
    next: Option<Rc<RefCell<Node>>>,  
}
```

Only RC<T> is not Enough

- RC<T> provides interior mutability via get_mut()
 - if counter = 1 (consider both shared and weak)
 - if cloned, get_mut() returns None during runtime

```
let mut x = Rc::new(1);  
let _w: Weak<_> = Rc::downgrade(&x);  
println!("a strong count: {:?}, weak count: {:?}",  
         Rc::strong_count(&x), Rc::weak_count(&x));  
let t1 = Rc::get_mut(&mut x).unwrap();  none  
/*t1 = 2;  
//assert_eq!(*x, 2);
```



RefCell<T>

- A mutable memory location.
- Perform borrow check during runtime.



```
let x = RefCell::new(1);  
{  
    let mut y = x.borrow_mut();  
    //let z = x.borrow_mut();  
    *y = 2;  
}  
assert_eq!(2, *x.borrow());
```

→ panic during runtime

RC vs RefCell



`get_mut()` may return `None` if cloned



the second `borrow_mut()` triggers panic

In-class Practice

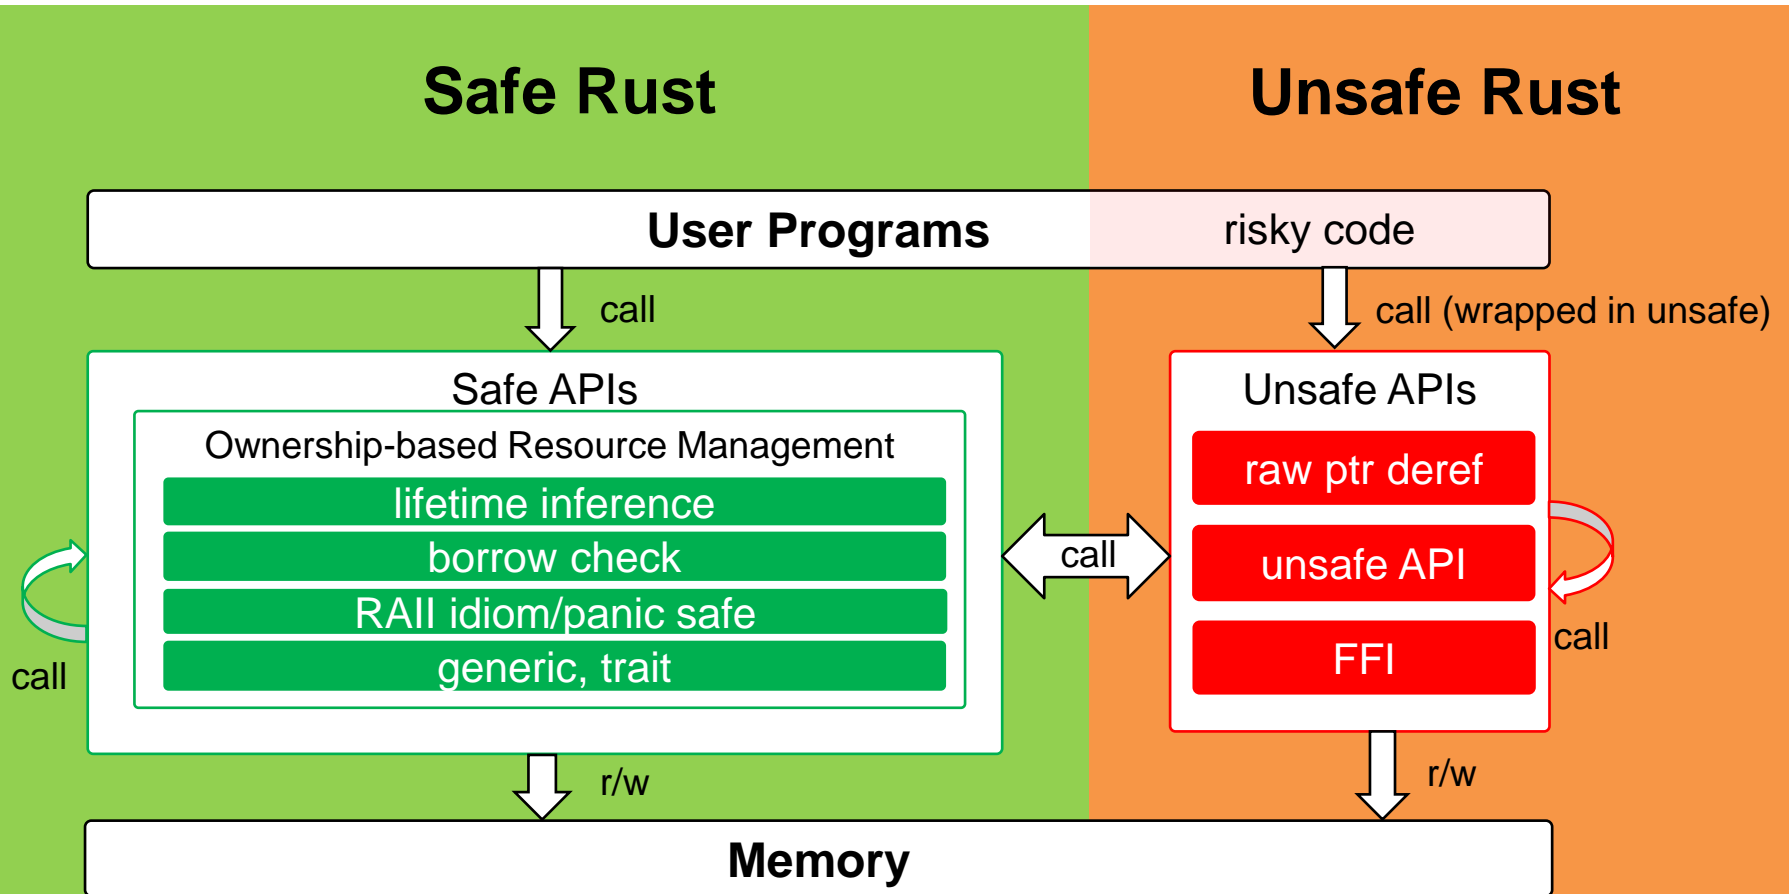
- Option 1: Implement a double linked list with Safe Rust
 - new
 - insert
 - delete
 - search
- Option 2: Implement a binary search tree with Safe Rust
 - new
 - insert
 - delete
 - search

Reference

- <https://doc.rust-lang.org/book/>
- <https://doc.rust-lang.org/stable/nomicon/>

Backup Slides

Rust Overview



Cell

- Perform borrow check during compile time;
- Check whether the content is Copy when calling get()

```
fn testcell(){  
    let mut x = Cell::new(1);  
    //let mut x = Cell::new(Box::new(1));  
    //x.set(2);  
    {  
        let mut y = x.get_mut();  
        //let z = x.get();  
        *y = 2;  
    }  
    assert_eq!(2, x.get());  
}
```

→ doesn't work

→ compile error

→ get() copies the value