

Lecture 7: Rust OBRM

徐 辉

xuh@fudan.edu.cn



Outline

- 1. Ownership
- 2. RAI and Lifetime
- 3. Unsafe Code

1. Ownership

Motivation of Design

- Dangling pointer is unacceptable
- Causal of dangling pointer?
 - Manual reclaim
 - we should prevent such manual reclaim
 - Automatic reclaim
 - alias analysis is NP-hard
 - shared pointer or gc is inefficient
- Rust comes to rescue

Overall Idea of Ownership

- Each object is owned by one variable.
 - Recall C++ `unique_ptr<T>`
- Ownership can be borrowed in two mode:
 - immutable: read only
 - mutable: read/write
- Exclusive mutability principle
 - two variables should not share mutable access to the same object at the same program point.
- Restrict the complexity of the alias analysis problem
 - we only need to trace the mutable pointer for each object
 - only one mutable pointer at each program point
 - only mutable pointers can lead to dangling pointers

Ownership & Borrowing

- Borrowed ownership will be returned automatically if no longer used.

```
fn main(){  
    let mut alice = Box::new(1);  
    let bob = alice;  
    println!("bob:{}", bob);  
    println!("alice:{}", alice);  
}
```

alice owns the object

transfer the vector to bob;
alice loses the ownership



```
fn main(){  
    let mut alice = Box::new(1);  
    let bob = &alice;  
    println!("bob:{}", bob);  
    println!("alice:{}", alice);  
}
```

bob borrows the ownership

bob returns the ownership to
alice automatically



Move Operator (=)

- If a type is not Copy (trait), move transfers the ownership.
 - e.g., Box<T> is not copy
- If a type is Copy, move does not transfer the ownership and only copies the value

```
fn main(){  
    let mut alice = 1;  
    let bob = alice;  
    println!("bob:{}", bob);  
    println!("alice:{}", alice);  
}
```

→ alice owns the object

→ copy the object to bob



Which Type Can be Copy?

- Primitive types on stack
- Composite types with all fields implementing copy
- How to (deep) copy objects of other non-Copy types:
 - implement Clone (trait)
 - each Copy type is also Clone.

```
fn main(){  
    let mut alice = Box::new(1);  
    let bob = alice.clone();  
    println!("bob:{}", bob);  
    println!("alice:{}", alice);  
}
```



→ alice owns the object

→ clone the object for bob

Mutability

mutable object

```
let mut alice = 1;  
alice+=1;
```



```
let alice = 1;  
alice+=1;
```



mutable borrow

```
let mut alice = 1;  
let bob = &mut alice;  
*bob+=1;
```



```
let mut alice = 1;  
let bob = &alice;  
*bob+=1;
```



```
let alice = 1;  
let bob = &mut alice;  
*bob+=1;
```



Mutability cont'd

```
let mut alice = 1;
```

```
let mut carol = 1;
```

```
let mut bob = &mut alice;
```

→ mutable object bob + mutable borrow

```
*bob+=1;
```

```
bob = &mut carol;
```

```
*bob+=1;
```



```
let mut alice = 1;
```

```
let mut carol = 1;
```

```
let bob = &mut alice;
```

→ immutable object bob + mutable borrow

```
*bob+=1;
```

```
bob = &mut carol;
```



→ bob cannot be modified

```
*bob+=1;
```

Debug Ownership Conflict

```
fn main(){
```

```
    let mut alice = 1;
```

```
    let bob = &mut alice;
```

```
    println!("bob:{}", bob);
```

```
    println!("alice:{}", alice);
```

```
}
```

mutable borrow

bob returns the ownership



```
fn main(){
```

```
    let mut alice = 1;
```

```
    let bob = &mut alice;
```

```
    println!("alice:{}", alice);
```

```
    println!("bob:{}", bob);
```

```
}
```


exclusive mutability



Rustviz Project: Visualize The Lifetime

Hover over timeline events (dots), states (vertical lines), and actions (arrows) for extra information.

```
1 fn main(){
2     let mut s = String::from("hello");
3
4     let r1 = &s;
5     let r2 = &s;
6     assert!(compare_strings(r1, r2));
7
8     let r3 = &mut s;
9     clear_string(r3);
10 }
```



Marcelo Almeida, Cyrus Omar, *et. al.*, "RustViz: Interactively Visualizing Ownership and Borrowing." In *2022 IEEE Symposium on Visual Languages and Human-Centric Computing (VL/HCC)*, pp. 1-10. 2022.

Pros and Cons

- Benefit
 - Compile-time prevention of shared mutable aliases
- When shared mutability is a Must...
 - Such as double linked lists?
 - Two options (we will discuss later):
 - shared pointer (reference counter)
 - unsafe code

2. RAI and Lifetime

RAI: Resource Acquisition is Initialization

Idea of RAI

- Ties resources to object lifetime
- Object/resource is initialized/allocated once created
 - all pointers refer to particular objects
 - no raw or dangling pointers
 - no uninitialized memory
- Resource deallocation is done during object destruction by the destructor
 - no manual deallocation is needed
 - achieved through static lifetime inference

Lifetime

- Each object has a lifetime constraint
- The object is reclaimed automatically after death
- A variable cannot borrow an object with a shorter lifetime

```
fn main(){  
    let alice;  
    {  
        let bob = 5;  
        alice = &bob;  
    }  
    println!("alice:{}", alice);  
}
```

→ bob lives in this subscope

→ alice points to an expired bob



Review Move for Non-Copy Types

- Extend the lifespan of the object on heap

```
fn test(){  
    let alice;  
    {  
        let bob = Box::new(1);  
        alice = bob;  
    }  
    println!("alice:{}", alice);  
}
```

→ transfer the ownership to alice



```
fn testret() -> Box<u64>{  
    Box::new(1)  
}
```

→ move the ownership to the ret value

```
let r = testret();  
println!("return:{}", r);
```



Lifetime Declaration

- Lifetime should be declared during function declaration

```
fn stringcmp(){
    let str1 = String::from("alice");
    let str2 = String::from("bob111");
    let result = longer(&str1, &str2);
    println!("The longer string is {}", result);
}

fn longer<'a>(x:&'a String, y:&'a String)->&'a String{
    if x.len()>y.len(){
        x
    } else {
        y
    }
}
```

Partial Order of Lifetime

- `<'a: 'b, 'b>` means a is relatively larger than b

```
fn stringcmp(){
    let str1 = String::from("alice");
    let result;
    //{
        let str2 = String::from("bob111");
        result = longer(&str1, &str2);
    //}
    println!("The longer string is {}", result);
}
```

```
fn longer<'a: 'b, 'b>(x:&'a String, y:&'b String)->&'b String{
    if x.len()>y.len(){
        x
    } else {
        y
    }
}
```

The return value cannot be 'a. Why?

Non-lexical Lifetime

- The default mode is non-lexical unless necessary
- Rust compiler tries to minimize the lifespan

```
'a: { let str1 = "alice";  
    'b: { let str2 = "bob";  
        'c: { let result = longer(str1,str2);  
              println!("The longer string is {}", result);  
        }  
    }  
}
```

Lifetime Elision to Be More Ergonomic

- Lifetime declaration can be elided if
 - there is only one input lifetime position
 - there are multiple positions, but one is `&self` or `&mut self`
 - assign the lifetime of `self` to elided output lifetimes

```
fn foo<'a>(s: &'a str, until: usize) -> &'a str;
```



```
fn foo(s: &str, until: usize) -> &str;
```



```
fn foo(s: &str, t: &str) -> &str; // ILLEGAL
```



```
fn foo<'a, 'b>(&'a mut self, t: &'b str) -> &str;
```



```
fn foo (&mut self, t: &str) -> &str;
```



More About Lifetime

- A static object means it lives for the entire lifetime
 - use the reserved lifetime 'static
 - all strings are static by default
- The lifetime can be unbounded
 - more flexible than static during lifetime inference

```
let s: &str = "hello world";
```



equivalent to

```
let s: &'static str = "hello world";
```

```
fn foo<'a>() -> &'a str;
```


unbounded lifetime

Automatic Reclaim/Drop

- Objects of Copy type (stack) can be reclaimed automatically
- For other objects with heap data?
 - Drop (trait) unused objects by calling the destructor
 - Recursively call the destructor of each field
- Drop and Copy are exclusive in Rust
 - Box<T> is Drop trait

```
struct MyType {a:u8, b:Box<u64>}
```

```
impl Drop for MyType {
```

```
    fn drop(&mut self){  automatically executed when the lifetime ends
```

```
        println!("dropping MyType object...");
```

```
    }
```

```
}
```

```
fn testdrop(){
```

```
    //create an object via "struct literal"
```

```
    let v = MyType {a:1, b:Box::new(2)};
```

```
}
```

Option<T> for Uninitialized Objects

- Option: an enumerate type
 - Some(T): the object type
 - None: if the object is uninitialized (null pointer)

```
pub enum Option<T> {  
    None,  
    Some(T),  
}  
  
let v = Some(...)  
match v.next {  
    Some(n) => ...,  
    None => panic!(),  
}
```


Example with a Singly-linked List

```
struct List{
    val: u64,
    next: Option<Box<List>>,
}

fn foo(){
    let mut l = List{val:1, next:None};
    l.next = Some(Box::new(List{val:2, next:None}));
    match l.next {
        Some(ref mut n) =>
            n.next = Some(Box::new(List{val:3, next:None})),
        None => panic!(),
    }
    let mut h = &l;
    loop {
        println!("{}", h.val);
        match h.next {
            Some(ref n) => h = n,
            None => break,
        }
    }
}
```

RAII for Thread Panic

- In a multi-threaded application, what happens when one thread exit exceptionally ?
 - abort: directly terminate the thread
 - panic: perform stack unwinding before exit
- Importance of RAII during stack unwinding
 - release locks (mutex)
 - release opened file descriptors
 - release allocated memories on heap

Sample Multi-thread Program


- When a spawned thread panics, unwind its stack
- The main thread continues execution
- Ineffective for fatal errors: e.g., stack overflow

```
fn main() {  
    let handle = thread::spawn(|| {  
        for i in 0..5 {  
            println!("new thread print {}", i);  
            thread::sleep(Duration::from_millis(10));  
        }  
        panic!();  
        //recursive();  
    });  
    for i in 0..10 {  
        println!("main thread print {}", i);  
        thread::sleep(Duration::from_millis(10));  
    }  
    handle.join();  
}
```

3. Unsafe Rust

Problem for Double-Linked List

- A node is owned by both its prev and next node
- Violate exclusive mutability



```
struct List{  
    val: u64,  
    prev: Option<Box<List>>,  
    next: Option<Box<List>>,  
}
```

Unsafe

- Operations that may lead to undefined behaviors are unsafe
 - Dereference raw pointers
 - Call unsafe functions
 - Call functions of foreign language (FFI)
- Using unsafe code within the unsafe scope

```
let mut num = 5;  
let r1 = &num as *const i32;  
println!("r1 is: {}", unsafe { *r1 });
```

→ r1 is a raw pointer

→ raw pointer dereference

Dereference raw pointers

```
unsafe fn foo() {  
    let addr = 0x012345usize;  
    let r = addr as *const i32;  
}  
unsafe { foo(); }
```


→ define an unsafe function

→ call the unsafe function

Call unsafe functions

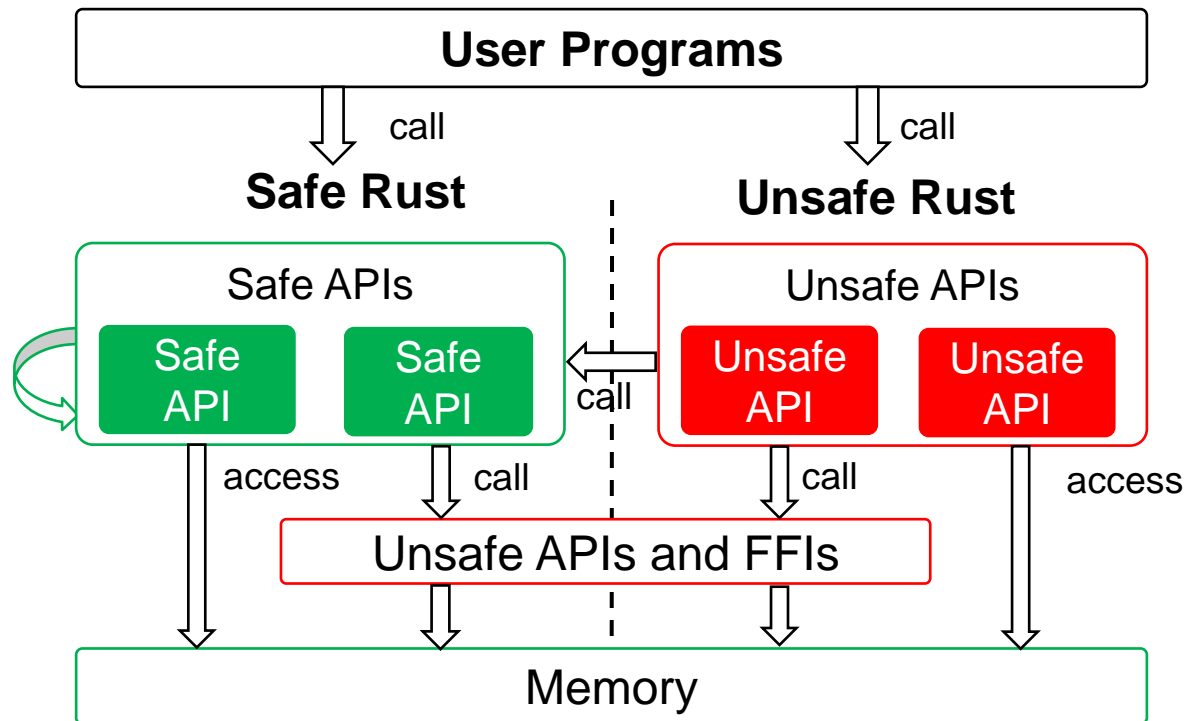
Trust Model

- Rust does not trust developers, so only safe code is allowed
- If a developer declares he knows the risk, Rust will trust him



Principle of Rust

- Safe API should not incur undefined behaviours
- Interior unsafe: wrap unsafe code into safe APIs
- Avoid using unsafe code unless necessary



Unsafe Version with Raw Pointers

- The objects pointed by raw pointers are not owned
- The resource may not be dropped automatically
- Prone to dangling pointers (Not RAII)

```
struct List{
    val: u64,
    next: *mut List,
    prev: *mut List,
}

fn foo(){
    let mut l = List{val:1, next:null_mut(), prev:null_mut()};
    l.next = &mut List{val:2, next:null_mut(), prev:null_mut()};
    unsafe {
        let mut cur = &mut *(l.next);
        cur.prev = &mut l;
        cur.next = &mut List{val:3, next:null_mut(), prev:null_mut()};
        (*(cur.next)).prev = cur;
    }
}
```

Solution Hint with RC<T> and RefCell<T>

- RC<T>: single-thread reference-counting pointer
 - enables shared immutable aliases
 - can mutate only if counter = 1 (cannot solve the list problem)
- RefCell<T>: a mutable memory location
 - convert immutable to mutable

```
struct List{  
    val: u64,  
    prev: Option<Rc<RefCell<List>>>,  
    next: Option<Rc<RefCell<List>>>,  
}
```

RC<T>

- Reference counter for shared aliases
- Mutate via get_mut()
 - mutual exclusion during compile time
 - if cloned, get_mut() returns None during run time

```
fn main(){  
    let mut x = Rc::new(1);  
    //let _y = Rc::clone(&x);  
    let t1 = Rc::get_mut(&mut x).unwrap();  
    //let t2 = Rc::get_mut(&mut x).unwrap();  
    *t1 = 2;  
    assert_eq!(*x, 2);  
  
    let _y = Rc::clone(&x);  
    assert!(Rc::get_mut(&mut x).is_none());  
}
```

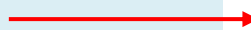
if cloned, get_mut()
returns None

compile error

RefCell<T>




- A mutable memory location
- Perform borrow check during runtime

```
fn testrefcell(){  
    let x = RefCell::new(Box::new(1));  
    {  
        let mut y = x.borrow_mut();  
        //let z = x.borrow_mut();  
        *(&y) = 2;  
    }  
    assert_eq!(2, *(&x.borrow()));  
}
```




panic during runtime

RC vs RefCell



get_mut() may return None if cloned



the second borrow_mut() triggers panic

In-class Practice

- 1) Implement a double linked list with Safe Rust
 - insert
 - delete
 - search
- 2) Implement a binary search tree with Safe Rust
 - insert
 - search