

Lecture 8: Rust Concurrency

Hui Xu

xuh@fudan.edu.cn



Outline

1. Multi-Thread Rust
2. Basic Concurrency APIs
3. Send/Sync

1. Multi-Thread Rust

Create a new Thread

```
use std::thread;
use std::time::Duration;
```

```
let tid = thread::spawn(|| {  
    println!("spawn");  
});
```

```
//thread::sleep(Duration::from_millis(1));
```

```
tid.join();
```

→ spawn a thread

→ wait for the thread ends

Create multiple Threads

```
let mut tids = vec![];  
  
for i in 0..3 {  
    tids.push(thread::spawn(move || {  
        ...  
    }));  
}  
  
for t in tids {  
    t.join();  
}
```

→ save thread handlers in a vector

Access Objects in a new Thread

- Access the same object from multiple threads is risky:
 - race condition
 - the thread may outlive the lifetime of the object

```
let mut x = 1;  
let _ = thread::spawn(move || {  
    x = 10;  
    println!("spawn: x = {}", x);  
}).join();  
println!("main: x= {}", x);
```

→ move the ownership or copy

→ copied x

```
spawn: x = 10  
main: x= 1
```

Access Objects of Drop Trait

```
let mut x = Box::new(1);  
let _ = thread::spawn(move || {  
    *x = 10;  
    println!("spawn: x = {}", x);  
}).join();  
println!("main: x= {}", x);
```



→ move the ownership of x to the thread

→ illegal to access x

```
let mut x = Box::new(1);  
let _ = thread::spawn(move || {  
    let mut y = x.clone();  
    *y = 10;  
    println!("spawn: y = {}", y);  
}).join();  
println!("main: x= {}", x);
```



→ make a clone of x as y

→ access y

```
spawn: y = 10  
main: x= 1
```

How to Share Data Sharing among Threads

```
let x = Box::new(1);  
let r = &x;  
let tid = thread::spawn(move || {  
    println!("{:?}", r);  
});
```



→ the thread may live longer than x

```
let x = Rc::new(Box::new(1));  
let cl = x.clone();  
let tid = thread::spawn(move || {  
    println!("{:?}", cl);  
});
```



→ RC is not thread safe

We Need Thread-Safe APIs

- Basic thread-safe APIs
 - Atomicity or lock
 - Synchronization or memory Barrier
- Advanced features


2. Basic Concurrency APIs

Atomic Types


- Several atomic types
 - AtomicBool,
 - AtomicIsize,
 - AtomicUsize,
 - ...
- Similar to C++ `std::atomic`

```
let mut foo = AtomicI32::new(0);
```

```
*foo.get_mut() = 5
```



```
foo.fetch_add(10, Ordering::SeqCst);
```



```
foo.compare_and_swap(5, 10, Ordering::Relaxed)
```




assignment

atomic add

CAS

Sample Mutex Lock with Memory Barrier

```
pub struct Mutex { flag: AtomicBool, }
impl Mutex {
    pub fn new() -> Mutex {
        Mutex { flag: AtomicBool::new(false), }
    }
    pub fn lock(&self) {
        while self.flag.compare_exchange_weak(
            false,                                // current
            true,                                // new
            Ordering::Relaxed, // success
            Ordering::Relaxed // failure
        ).is_err() {}
        fence(Ordering::Acquire);
    }
    pub fn unlock(&self) {
        self.flag.store(false, Ordering::Release);
    }
}
```



all subsequent loads will see the stored data

Mutex provided in the Standard Library

- `lock()` is blocking mode that blocks the thread until successful
- `try_lock()` is nonblocking mode
 - returns `Err()` if fails

```
let x = Mutex::new(0);  
let _ = thread::spawn(move || {  
    let mut data = x.lock().unwrap();  
    *data += 1;  
    println!("{:?}", data);  
}).join();
```

→ Do not need to unlock, why?

Arc<T>: Atomically Ref Counted

- Mutex is not enough; it cannot be shared among threads.
- Share ownership through Arc.
 - Similar to Rc<T>, but thread safe.
- Use atomic operations for reference counting

```
let x = Arc::new(Mutex::new(0));
let cl = x.clone();
let tid = thread::spawn(move || {
    let mut data = cl.lock().unwrap();
    *data += 1;
    println!("{:?}", data);
});
tid.join();
let y = x.lock().unwrap();
println!("{:?}", y);
```



What if moving this line of code to the end?

Mutex: Poison Strategy

- What if a thread holding the lock panics?
- Using a poison flag to detect/recover from the bad state

```
let arc = Arc::new(Mutex::new(0));  
let cl = arc.clone();
```

```
let _ = thread::spawn(move || -> () {  
    let mut data = cl.lock().unwrap();  
    panic!();  
}).join();
```

→ Panic the thread

```
assert_eq!(arc.is_poisoned(), true);  
let mut guard = match arc.lock() {  
    Ok(guard) => guard,  
    Err(poisoned) => poisoned.into_inner(),  
};  
*guard += 1;
```

→ The lock is poisoned

→ Release the locked data

Synchronizing Primitive: Condition Variable

- Do not consume CPU when threads need to wait for a resource to become available.
- How to implement the feature? (OS hangs up the thread)

```
let x = Arc::new((Mutex::new(0), Condvar::new()));
let cl = Arc::clone(&x);
let tid = thread::spawn(move || {
    let (l, c) = &*cl;
    let mut t = l.lock().unwrap();
    *t = 100;
    println!("spawned thread, t = {}", t);
    c.notify_one();
});
let (l, c) = &*x;
let mut t = l.lock().unwrap();
t = c.wait(t).unwrap(); // release the original lock
println!("main thread, t = {}", t);
tid.join();
```

```
spawned thread, t = 100
main thread, t = 100
```


Effectiveness of Condition Variable

```
let x = Arc::new((Mutex::new(0), Condvar::new()));
let cl = Arc::clone(&x);
let tid = thread::spawn(move || {
    let (l, c) = &*cl;
    let mut t = l.lock().unwrap();
    *t = 100;
    println!("spawned thread, t = {}", t);
    c.notify_one();
});
let (l, c) = &*x;
let mut t = l.lock().unwrap();
*t = 2;
println!("main thread, t = {}", t);
t = c.wait(t).unwrap();
println!("main thread, t = {}", t);
tid.join();
```

```
main thread, t = 2
spawned thread, t = 1
main thread, t = 1
```

More: Message Passing

- Multi-producer, single-consumer FIFO queue
 - Asynchronous or synchronous mode

```
let (tx, rx) = mpsc::channel();
let tx = tx.clone();
let _ = thread::spawn(move || {
    for i in 0..10 {
        tx.send(i).unwrap();
    }
}).join();

while let Ok(msg) = rx.recv(){
    println!("receive: {}",msg);
}
```


3. Send/Sync

Marker Traits

- Marker Traits have no methods to implement
- They are compiler intrinsic and auto derived
 - Send/!Send
 - Sync/!Sync
- Other marker traits
 - Copy/!Copy
 - Sized/!Sized
 - Unpin/!Unpin


Send

- The type of Send can be transferred between threads.
- Use the move operator, which is similar as =
 - For types of Copy trait, make a copy of the object.
 - For types of Drop trait, transfer the ownership.
- Almost all primitive types are Send.
- Any struct composed of Send types is automatically marked as Send.



Sync

- The type of Sync is safe to be referenced from multiple threads.
- Any type T is Sync if $\&T$ is Send.
- Sync is usually more rigid than Send. Why?



Raw pointers are neither Send nor Sync

- Possible to create shared objects (although unsafe).
- Should be manually implemented as unsafe.



```
struct Unsend{ ptr: *mut i64, }  
impl Unsend{  
    fn add(&self, i:i64){  
        unsafe{*(self.ptr) = *self.ptr + i};  
    }  
}  
unsafe impl Send for Unsend{ }  
unsafe impl Sync for Unsend{ }
```

```
let mut var = 0i64;  
let mut v = Unsend{ptr:&mut var as *mut i64};  
let tid = thread::spawn(move || {  
    for i in 1..100001{ v.add(i); }  
});  
for i in 1..100001{ var+=i; }  
tid.join();  
println!("{}",var);
```

→ Implement Send/Sync is unsafe

Can Cell/RefCell Be Send/Sync?

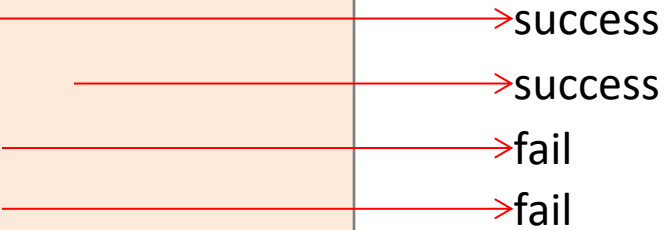
- For unsynchronized interior mutability.
- They are Send but not Sync.



Code to Verify Send/Sync Properties

```
fn check_send<T: Send>(_: T) {}
fn check_sync<T: Sync>(_: T) {}

fn testCell(){
    let mut v = Cell::new(1);
    let mut v = RefCell::new(1);
    check_send(v);           → success
    check_send(&mut v);      → success
    //check_send(&v);         → fail
    //check_sync(v);         → fail
}
```



Rc<T> and Arc<T>

- Rc<T> is neither Send nor Sync, why?
 - !Send: cloned Rc exist in multiple threads
 - atomicity in reference counter update
- Does Arc<T> have bound on T with Sync to be thread-safe?
 - The compiler checks the wrapped data during compilation

```
impl<T> !Send for Rc<T>
impl<T> !Sync for Rc<T>
```

→ Force Rc<T> to be !Send + !Sync

```
let mut v = Arc::new(Cell::new(1));
let v1 = v.clone();
thread::spawn(move || {
    (*v1).set(3);
}).join();
(*v).set(2);
```

→ No manual bound on T

→ Compilation error

Can Mutex be Send/Sync?

- Require T to be Send

```
let mut v = Mutex::new(1);  
check_send(v);           _____> success  
check_send(&v);           _____> success  
check_sync(v);           _____> success
```

```
let mut v = Mutex::new(Cell::new(1));  
check_send(v);           _____> success  
check_send(&v);           _____> success  
check_send(&mut v);       _____> success  
check_sync(v);           _____> success
```

Can Mutex be Send/Sync? Cont'd

```
Let mut cell = Cell::new(1)
let mut v = Mutex::new(&cell);
check_send(v);           _____>fail
//check_sync(v);         _____>fail
```

```
Let mut cell = Cell::new(1)
let mut v = Mutex::new(&mut cell);
check_send(v);           _____>success
check_sync(v);           _____>success
```

Sync but not Send?

- Cases are rare.
 - thread-local features, *e.g.*, MutexGuard

In-Class Practice

- Rewrite your program (binary search tree) to be thread-safe.
 - Support Sync/Send
- Discuss why your program is thread safe.
- Design experiments to show that your program is thread safe.

Backup Slides

Synchronizing Primitive: Once

- Run global initialization only one time
 - access 'static mut' variables

```
static mut VAL: usize = 0;
static INIT: Once = Once::new();

fn get_cached_val() -> usize {
    unsafe {
        INIT.call_once(|| {
            VAL = expensive_computation();
        });
        VAL
    }
}
```