

# Lecture 9: Rust Concurrency

徐辉

xuh@fudan.edu.cn



# Outline

1. Multi-Thread Rust
2. Basic Concurrency APIs
3. Marker Trait for Concurrency

# 1. Multi-Thread Rust

---

# Spawn A New Thread

```
use std::thread;
use std::time::Duration;
```

```
fn main() {
    let thd = thread::spawn(|| {
        ...
        thread::sleep(Duration::from_millis(1));
    });
    thd.join().unwrap();
}
```

spawn a thread

wait for the thread  
ends

# Spawn A Group of Threads

```
let mut threads = vec![];  
  
for i in 0..3 {  
    threads.push(thread::spawn(move || {  
        ...  
    }));  
}  
  
for t in threads {  
    t.join();  
}
```

save thread handlers in a vector

# Access Objects In the New Thread

- Access the same object from multiple threads is risky:
  - race condition
  - the thread may outlive the lifetime of the object

```
let mut x = 1;
//let tid = thread::spawn(|| {
let tid = thread::spawn(move || {
    x = 10;
    println!("spawn: x = {}", x);
});
tid.join().unwrap();
println!("main: x= {}", x);
```

— move the ownership or copy

— copied x

```
spawn: x = 10
main: x= 1
```

# Object Access: Cont'd

```
let mut x = Box::new(1);
let tid = thread::spawn(move || {
    *x = 10;
    println!("spawn: x = {}", x);
});
tid.join().unwrap();
println!("main: x= {}", x);
```

— move the ownership of x to the thread

✗ — Illegal to access x

```
let mut x = Box::new(1);
let mut y = x.clone();
let tid = thread::spawn(move || {
    *y = 10;
    println!("spawn: y = {}", y);
});
tid.join().unwrap();
println!("main: x= {}", x);
```

— make a clone of x as y

— access y

✓

```
spawn: y = 10
main: x= 1
```

# Share Data Among Threads

```
let x = Box::new(1);
for i in 0..3 {
    let r = &x;
    thread::spawn(move || {
        println!("{:?}", r);
    });
}
```



the thread may live longer than x

```
let x = RC::new(Box::new(1));
for i in 0..3 {
    let cl = x.clone();
    thread::spawn(move || {
        println!("{:?}", cl);
    });
}
```



RC is not thread safe



# We Need Concurrency-Safety APIs

- Basic APIs
  - Atomicity or lock
  - Synchronization or memory Barrier
- Advanced features

## 2. Basic Concurrency APIs

---

# Atomic Types

- Several atomic types
  - AtomicBool,
  - AtomicIsize,
  - AtomicUsize,
  - ...
- Similar to C++ `std::atomic`

```
let mut foo = AtomicI32::new(0);  
*foo.get_mut() = 5  
foo.fetch_add(10, Ordering::SeqCst);  
foo.compare_and_swap(5, 10, Ordering::Relaxed);
```

assignment  
atomic add  
CAS

# Sample Mutex Lock with Memory Barrier

```
pub struct Mutex { flag: AtomicBool, }  
impl Mutex {  
    pub fn new() -> Mutex {  
        Mutex { flag: AtomicBool::new(false), }  
    }  
    pub fn lock(&self) {  
        while self.flag.compare_exchange_weak(  
            false, true,  
            Ordering::Relaxed, Ordering::Relaxed)  
                .is_err() {}  
        fence(Ordering::Acquire);  
    }  
    pub fn unlock(&self) {  
        self.flag.store(false, Ordering::Release);  
    }  
}
```

———— all subsequent loads will see the stored data

# Arc<T>: Atomically Ref Counted

- Similar to RC<T>, but is thread safe
- Use atomic operations for reference counting
- Mutating through an Arc generally use Mutex, RwLock, etc.

```
let x = Arc::new(Box::new(1));  
for i in 0..3 {  
    let cl = x.clone();  
    thread::spawn(move || {  
        println!("{:?}", cl);  
    });  
}
```



# Mutex

- Use `lock()` or `try_lock()` to access the data
  - Returns `Result<T>`
  - `lock()` is blocking mode
    - most usage simply `unwrap()` the result, why?
  - `try_lock()` is nonblocking mode
    - returns `Err()` if fails

```
let x = Arc::new(Mutex::new(0));
```

```
for i in 0..3 {  
    let cl = x.clone();  
    thread::spawn(move || {  
        let mut data = cl.lock().unwrap();  
        *data += 1;  
        println!("{:?}", data);  
    });  
}
```

Do not need to unlock, why?

# Synchronizing Primitive: Condition Variable

- Do not consume CPU when threads need to wait for a resource to become available
- How to implement the feature?

```
let x = Arc::new((Mutex::new(0), Condvar::new()));
let cl = Arc::clone(&x);
thread::spawn(move || {
    let (l, c) = &*cl;
    let mut t = l.lock().unwrap();
    *t = 100;
    c.notify_one();
});
let (l, c) = &*x;
let mut t = l.lock().unwrap();
while *t == 0 {
    t = c.wait(t).unwrap();
    println!("while: t = {}", t);
}
```

# Mutex: Poison Strategy

- What if a thread holding the lock panics?
- Using a poison flag to detect/recover from the bad state

```
let arc = Arc::new(Mutex::new(0));  
let cl = arc.clone();
```

```
let _ = thread::spawn(move || -> () {  
    let mut data = cl.lock().unwrap();  
    panic!();  
}).join();
```

— Panic the thread

```
assert_eq!(arc.is_poisoned(), true);
```

— The lock is poisoned

```
let mut guard = match arc.lock() {  
    Ok(guard) => guard,  
    Err(poisoned) => poisoned.into_inner(),  
};
```

— Release the locked data

```
*guard += 1;
```



# Message Passing

- Multi-producer, single-consumer FIFO queue
  - Asynchronous or synchronous mode

```
use std::sync::mpsc;
use std::thread;
fn main() {
    let (tx, rx) = mpsc::channel();
    let tx = tx.clone();
    let tid = thread::spawn(move || {
        for i in 0..10 {
            tx.send(i).unwrap();
        }
    });

    while let Ok(msg) = rx.recv(){
        println!("receive: {}",msg);
    }
}
```

# Synchronizing Primitive: Once

- Run global initialization only one time
  - access 'static mut' variables

```
static mut VAL: usize = 0;
static INIT: Once = Once::new();

fn get_cached_val() -> usize {
    unsafe {
        INIT.call_once(|| {
            VAL = expensive_computation();
        });
        VAL
    }
}
```

### 3. Marker Trait for Concurrency


---

# Marker Traits for Concurrency

- Marker Traits have no methods to implement
- They are compiler intrinsic and auto derived
  - Send/!Send
  - Sync/!Sync
- Other marker traits
  - Copy/!Copy
  - Sized/!Sized
  - Unpin/!Unpin


# Send

- The type of Send can be transferred between threads
- Use the move operator, which is similar as =
  - For types of Copy trait, make a copy of the object
  - For non-copy, transfer the ownership
- Almost all primitive types are Send
- Any struct composed of Send types is automatically marked as Send



# Sync

- The type of Sync is safe to be referenced from multiple threads
- Any type T is Sync if &T is Send
- Sync is usually more rigid than Send. Why?



# Raw pointers are neither Send nor Sync

- Possible to create shared objects (although unsafe)
- Should be manually implemented as unsafe

```
struct Unsend{ ptr: *mut i64, }  
impl Unsend{  
    fn add(&self, i:i64){  
        unsafe{*(self.ptr) = *self.ptr + i};  
    }  
}
```



```
unsafe impl Send for Unsend{  
unsafe impl Sync for Unsend{
```

Implement Send/Sync is unsafe

```
let mut var = 0i64;  
let mut v = Unsend{ptr:&mut var as *mut i64};  
let tid = thread::spawn(move || {  
    for i in 1..100001{ v.add(i); }  
});  
for i in 1..100001{ var+=i; }  
tid.join();  
println!("{}",var);
```

# Can Cell/RefCell Be Send/Sync?

- unsynchronized interior mutability
- Send but not Sync





# Code to Verify Send/Sync Properties

```
fn check_send<T: Send>(_: T) {}
fn check_sync<T: Sync>(_: T) {}

fn testCell(){
  let mut v = Cell::new(1);
  let mut v = RefCell::new(1);
  check_send(v);           success
  //check_send(&mut v);    success
  //check_send(&v);         fail
  //check_sync(v);         fail
}
```

# Rc<T> and Arc<T>

- Rc<T> is neither Send nor Sync, why?
  - !Sync: atomicity in reference counter update
  - !Send: cloned Rc exist in multiple threads
- Does Arc<T> have bound on T to be thread-safe?
  - The compiler checks the wrapped data during compilation

```
impl<T> !Send for Rc<T> ————— Force Rc<T> to be !Send + !Sync
impl<T> !Sync for Rc<T>
```

```
let mut v = Arc::new(Cell::new(1)); ————— No Bound on T
let v1 = v.clone();
thread::spawn(move || { ————— Compilation error
    (*v1).set(3);
}).join();
(*v).set(2);
```

# Can Mutex be Send/Sync?

- Require T is Send

```
let mut v = Mutex::new(1);  
//check_send(v);           _____ success  
//check_send(&v);          _____ success  
check_sync(v);             _____ success
```

```
let mut v = Mutex::new(Cell::new(1));  
check_send(v);             _____ success  
//check_send(&v);          _____ success  
//check_send(&mut v);      _____ success  
//check_sync(v);           _____ success
```

# Can Mutex be Send/Sync? Cont'd

```
let mut v = Mutex::new(&Cell::new(1));  
check_send(v);  
//check_sync(v);
```

fail  
fail

```
let mut cell = Cell::new(1)  
let mut v = Mutex::new(&mut cell);  
check_send(v);  
//check_sync(v);
```

success  
success

# Sync but not Send?

- Cases are rare
- Exceptions may relate to thread-local features,
  - e.g., MutexGuard

# In-Class Practice

- Rewrite your program (binary search tree or double-linked list) to be thread-safe
  - Support Sync/Send
- Show that your program is thread safe